From Content to Action: Agentic AI in IBM Content Cortex

Instructor — led lab guide

Malek Jabri
AI Assistant/ Business Automation EMEA
Technical Sales
Malek.jabri@be.ibm.com

Table of Contents

1 Introduction 6

1.1 From document assistance to agentic content management 6

1.2 What you'll learn 6

1.3 Lab journey and outcomes 7

1.4 Lab context : 7

1.4.1 Document Class 8

1.4.2 Document Properties (The "Metadata" or "Attributes") 8

2 Prerequisites 9

2.1 Test your credentials 9

3 Lab Exercise 1: Grounded document insight with Content Assistant 11

3.1 Ask, verify, and refine 11

3.2 Compare different documents 15

4 Lab Exercise 2: Navigator Core Agents 19

4.1 Introduction 19

4.2 Technical overview 20

4.3 Tour of the new features 21

4.3.1 Let's switch to Ask AI Agents: 21

4.4 IBM Navigator: Ask AI Agents 22

4.4.1 Welcome Screen 25

4.4.2 Find documents 26

4.4.3 Update documents 32

5 Lab Exercise 3: Integrate Content Cortex capabilities 41

5.1 Content Cortex Agentic Integration: 41

5.1.1 Let's get the lab starter : 41

- Change the username to the one assigned to you by the instructor. 42

5.2 Content Cortex Assistant Integration (Optional) : 50

5.2.1 Graphql 51

5.2.2 Lab Setup Instructions 52

5.2.3 Ask a generic question. 54

5.2.4 Ask your repository 57
5.2.5 Ask a question on a document. 58
5.2.6 More exercises: 61

6 Lab Exercise 5: Configure and govern AI-enabled content (Optional) 62

6.1 Introduction 62

6.2 Enable your content assistant and samples (Front) 62

6.3 Administration and operational assurance 67
6.3.1 Review Content Cortex backbone configuration 67
6.3.2 Enhance Extraction (New feature) 71
6.3.3 Monitor AI requests and inspect traceability 71

Notices and disclaimers

© 2025 International Business Machines Corporation. No part of this document may be reproduced or transmitted in any form without written permission from IBM.

U.S. Government Users Restricted Rights — use, duplication or disclosure restricted by GSA ADP Schedule Contract with IBM.

This document is current as of the initial date of publication and may be changed by IBM at any time. Not all offerings are available in every country in which IBM operates.

Information in these presentations (including information relating to products that have not yet been announced by IBM) has been reviewed for accuracy as of the date of initial publication and could include unintentional technical or typographical errors. IBM shall have no responsibility to update this information.

This document is distributed “as is” without any warranty, either express or implied. In no event, shall IBM be liable for any damage arising from the use of this information, including but not limited to, loss of data, business interruption, loss of profit or loss of opportunity. IBM products and services are warranted per the terms and conditions of the agreements under which they are provided. The performance data and client examples cited are presented for illustrative purposes only. Actual performance results may vary depending on specific configurations and operating conditions.

IBM products are manufactured from new parts or new and used parts.
In some cases, a product may not be new and may have been previously installed. Regardless, our warranty terms apply.”

Any statements regarding IBM's future direction, intent or product plans are subject to change or withdrawal without notice.

Performance data contained herein was generally obtained in a controlled,
isolated environments. Customer examples are presented as illustrations of how those customers have used IBM products and the results they may have achieved. Actual performance, cost, savings or other results in other operating environments may vary.

References in this document to IBM products, programs, or services does not imply that IBM intends to make such products, programs or services available in all countries in which IBM operates or does business.

Workshops, sessions and associated materials may have been prepared by independent session speakers, and do not necessarily reflect the views of IBM. All materials and discussions are provided for informational purposes only, and are neither intended to, nor shall constitute legal or other guidance or advice to any individual participant or their specific situation.

It is the customer’s responsibility to ensure its own compliance with legal requirements and to obtain advice of competent legal counsel as to the identification and interpretation of any relevant laws and regulatory requirements that may affect the customer’s business and any actions the customer may need to take to comply with such laws. IBM does not provide legal advice or represent or warrant that its services or products will ensure that the customer follows any law.

Notices and disclaimers (Continued)

Questions on the capabilities of non-IBM products should be addressed to the suppliers of those products. IBM does not warrant the quality of any third-party products, or the ability of any such third-party products to interoperate with IBM's products. IBM expressly disclaims all warranties, expressed or implied, including but not limited to, the implied warranties of merchantability and fitness for a purpose.

The provision of the information contained herein is not intended to, and does not, grant any right or license under any IBM patents, copyrights, trademarks or other intellectual property right.

IBM, the IBM logo, and ibm.com are trademarks of International Business Machines Corporation, registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the Web at "Copyright and trademark information" at: www.ibm.com/legal/copytrade.shtml.

1 Introduction

1.1 From document assistance to agentic content management

Enterprise content is one of the largest sources of trusted business context, but it is often fragmented, difficult to interpret, and disconnected from automation. **IBM Content Cortex** is an agentic content services platform that centralises, governs, and activates content so people and AI agents can work with it securely at enterprise scale.

This hands-on lab introduces the Content Cortex journey from governed repository content to grounded insight and agentic action. You will use **IBM Content Assistant** for document-level generative AI, explore semantic retrieval and repository-wide questions, compare content, invoke capabilities through APIs, and examine the governance and operational controls that make enterprise AI trustworthy.

1.2 What you'll learn

By completing the lab, you will be able to:

Explain the Content Cortex value proposition — Describe how governed repositories, content services, semantic retrieval, reasoning, conversational experiences, and standards-based integration work together.

- Use selected documents, folders, search results, and the document viewer as grounded scopes.
- Ask repository-wide questions and inspect supporting sources.
- Compare documents and versions to identify meaningful changes.
- Distinguish insight generation from agentic content operations.

Apply Content Assistant — Ask natural-language questions, generate summaries, retrieve structured facts, and compare content by using context-specific Retrieval-Augmented Generation (RAG).

Explore agentic integration and governance — Understand how Content Cortex AI Services, including conversational agents, reasoning, and Model Context Protocol (MCP) interfaces, can expose secure content operations to AI applications while repository permissions and lifecycle controls remain in force.

Configure prompts, models, indexing scope, and eligible document classes.

Review source grounding, query history, and operational records.

Use GraphQL to embed content intelligence in custom applications.

Design a controlled agentic workflow that moves from finding information to proposing and validating an action.

1.3 Lab journey and outcomes

You will work with a preconfigured, non-sensitive content set and progress through six connected exercises: grounded document insight, repository discovery, document and version comparison, API-based integration, administration and monitoring, and an agentic content scenario. Each exercise includes an objective, guided tasks, validation points, and a short reflection.

By the end, you will be able to explain where Content Assistant fits within Content Cortex, demonstrate grounded answers and source inspection, integrate generative AI with content services, and outline how governance, human oversight, and auditability support safe agentic operations.

1.4 Lab context :

IBK is a fictional global financial institution with six lines of business: Retail Banking, Corporate & Institutional Banking, Wealth Management, Insurance, Compliance & Risk, and Operations & Back-Office. The data model provides the repository foundation — document classes, folder classes, property templates, and choice lists — required to run scripted agentic demonstrations showcasing AI-driven document capture, extraction, classification, and lifecycle management across all divisions.

The model is structured around a shared base class (IBKDocument / IBKFolder) that carries common metadata applicable to every document in the organisation. Each line of business extends this base with its own specialised properties and folder hierarchy, enabling both unified enterprise-wide search and fine-grained, context-aware classification per division.

All symbolic names use the IBK prefix and English-only Latin characters to ensure long-term stability in FileNet. Sample document content spans English, German, Spanish, and French to demonstrate the agentic layer's cross-language extraction capabilities.

We have defined the document and folder type topology:

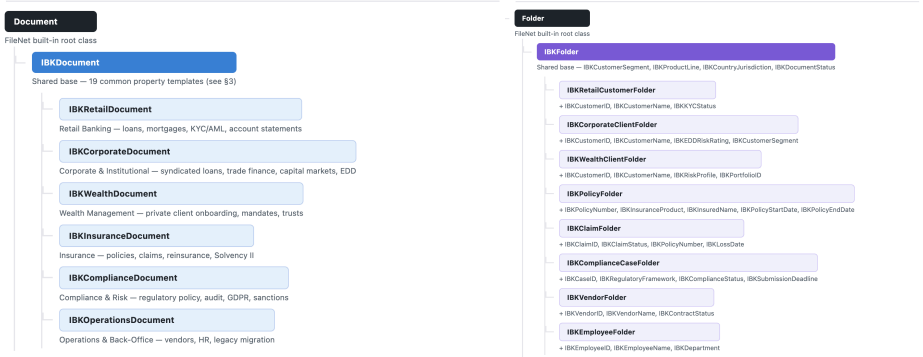

1.4.1 Document Class

A **Document Class** defines the **type or category** of a business record. It acts as a standardized blueprint that dictates what kind of document it is, what data it must capture, who can access it, and how long it must be kept.

- **Business Definition:** A grouping mechanism that organises documents sharing the same business purpose, regulatory rules, and data structure.

- **Key Business Roles:**
  - **Governance & Retention:** Governs how long the document is retained (e.g., tax records kept for 7 years vs. employment contracts kept for duration of employment + 5 years).
  - **Security & Access Control:** Controls which business units or roles can view, edit, or approve the document.
  - **Lifecycle & Workflow:** Triggers specific business processes (e.g., assigning an IBKRetailDocument directly to a mortgage approval workflow).

- **Examples from Banking/Finance:**
  - MortgageApplication
  - CommercialLoanAgreement
  - ComplianceReport (SAR)
  - EmploymentContract

1.4.2 Document Properties (The "Metadata" or "Attributes")

Properties (also called **Metadata**, **Attributes**, or **Index Fields**) are the **specific business data points** attached to an individual document to describe its context, content, and state without needing to open the file.

- **Business Definition:** Key-value data elements associated with a document that make it searchable, auditable, and actionable within business workflows.

- **Key Business Roles:**
  - **Searchability & Retrieval:** Allows business users to locate documents instantly (e.g., searching by *Customer ID* or *Policy Number*).
  - **Process Automation:** Feeds data into downstream core banking, ERP, or CRM systems.
  - **Audit & Compliance:** Tracks critical identifiers, creation dates, and involved counterparties.

- **Examples for a** MortgageApplication **Class:**
  - **Customer ID / SSN:** CUST-849201
  - **Applicant Name:** Maria Garcia
  - **Loan Amount:** $450,000
  - **Submission Date:** 2024-03-15
  - **Status:** Pending Review
  - **Branch / Region:** North America - Retail

2 Prerequisites

You should have received either:

- An email with the different links and credentials
- The Box link containing the cluster access information.

  https://ibm.box.com/v/Content-Cortex-Lab

2.1 Test your credentials

You should have received your credentials to access the server, along with the lab documentation and resources needed to complete the exercises.

Open the URL:

You will be redirected to the following authentication portal.

From the “Log in with” drop-down list, select “Enterprise LDAP.”

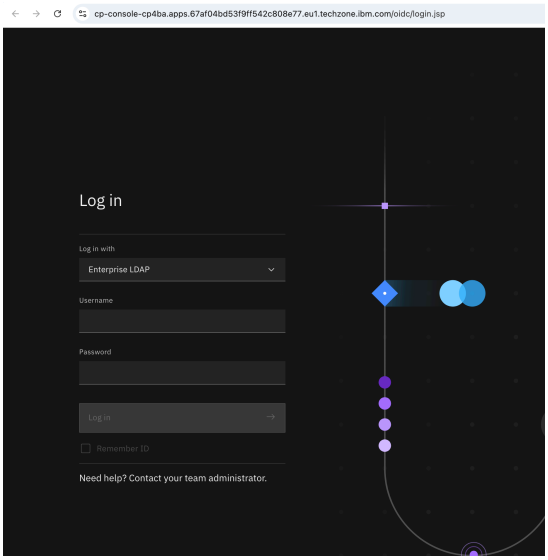

After logging in, you will be redirected to the Navigator page.

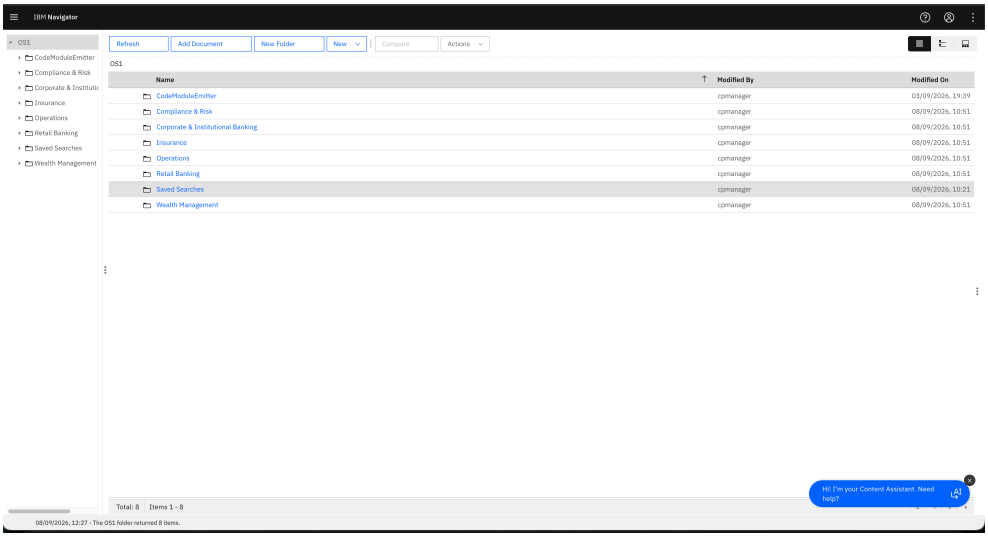

Nice, you are now connected to Navigator.

3 Lab Exercise 1: Grounded document insight with Content Assistant

3.1 Ask, verify, and refine

Objective: Use Content Assistant across single-document, multi-document, and folder scopes; validate every important answer against its cited source; and observe how changing scope changes the response.

Navigate to the following folder: OS1/Compliance & Risk/AML & Sanctions

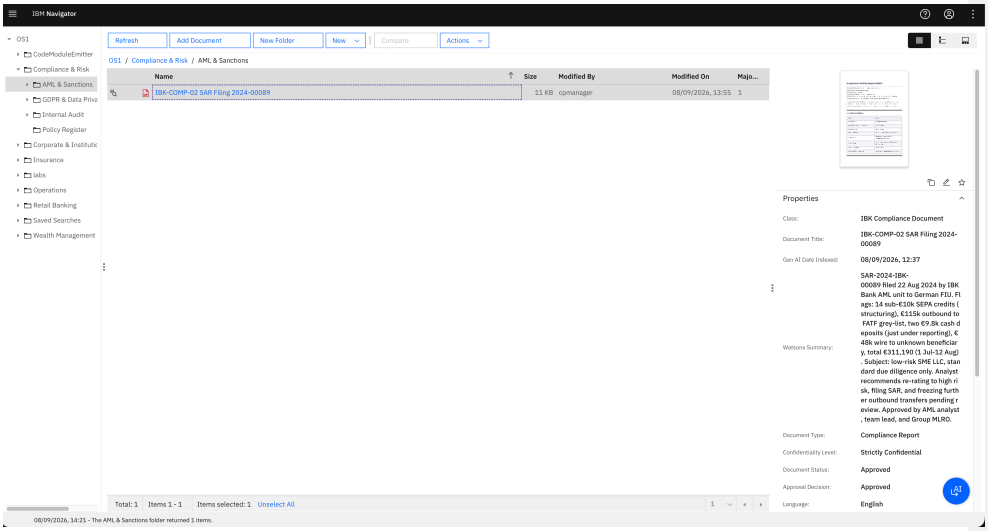

You will find a single document in the folder. An index icon appears next to the document icon, indicating the document has been indexed in the vector database for retrieval-augmented generation (RAG).

Double-click the first document in the list to review its content using the embedded viewer. After reviewing the document, click the X button in the top-right corner of the viewer to close it.

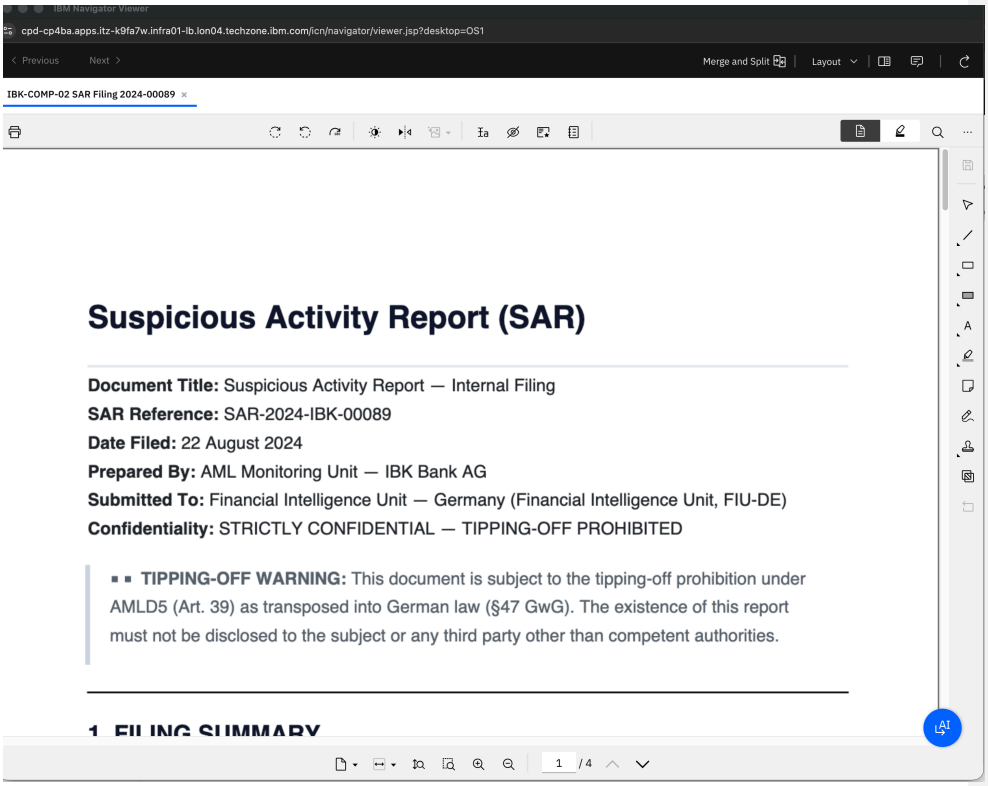

The Content Assistant generates a summary for each document. Select the first document in the list, and a panel on the right side will be populated with a thumbnail and a set of key properties, including a summary.

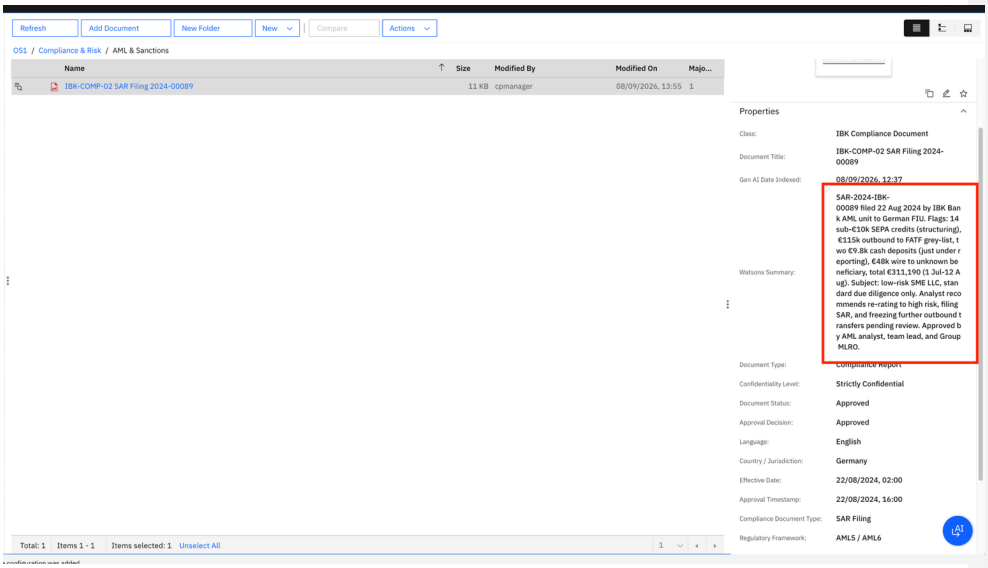

Click on the blue icon at the bottom of the page to launch the assistant.

At the bottom of the chat, you will see the message “1 document selected.”

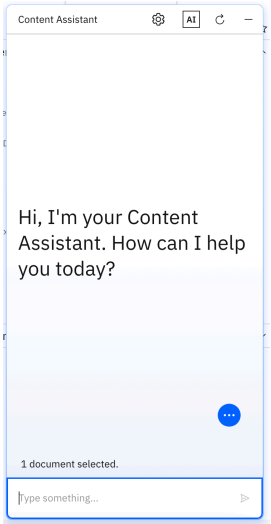

This means that the assistant will answer questions using only the content of that document. For example, you can ask the following questions:

Ask: “What is the SAR Reference, and what are the suspicious activities?”

Ask: “Who are the analysts?”

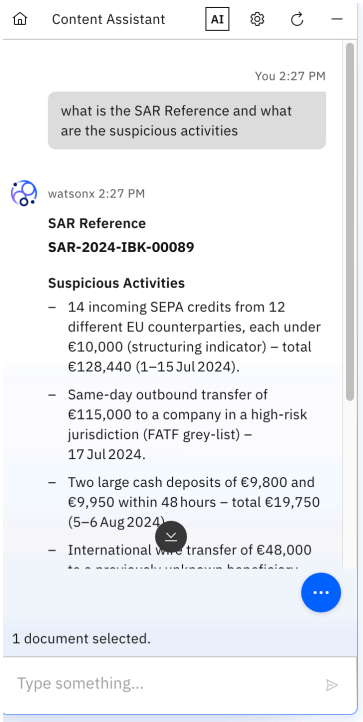

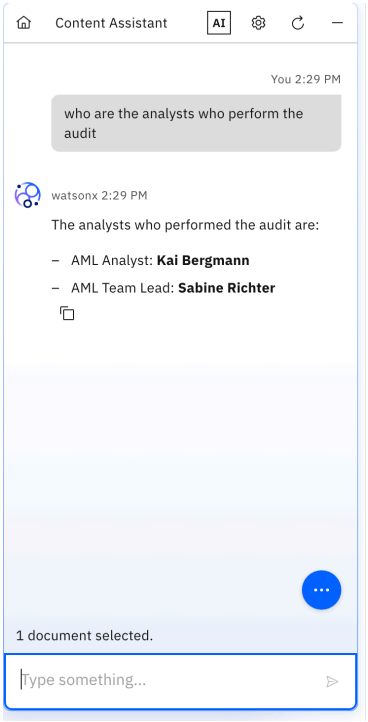

Ask: Could you provide me with a summary.

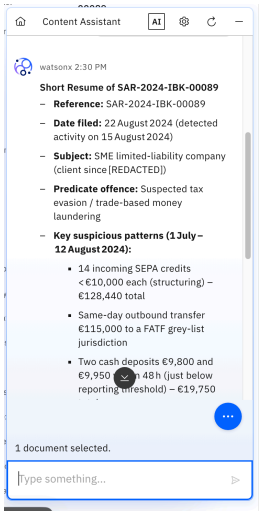

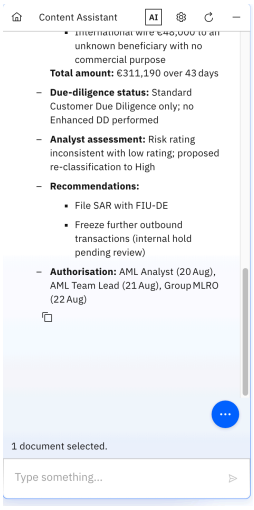

Validation checkpoint: For one extracted fact, open the supporting source and confirm the value in the original document. Record whether the answer is exact, partially supported, or unsupported. If unsupported, refine the question or reduce the scope rather than accepting the response.

You can rate the answer by selecting a different option or copying the response.

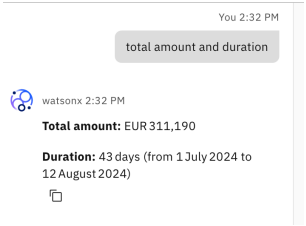

Objective: Use generative AI to compare related documents and successive versions, then validate the material differences against source text. Focus on changed obligations, coverage, fees, dates, limits, exclusions, and other business-significant terms rather than cosmetic wording.

3.2 Compare different documents

- Navigate to the folder: OS1/Insurance/Life Insurance — Karl Fischer

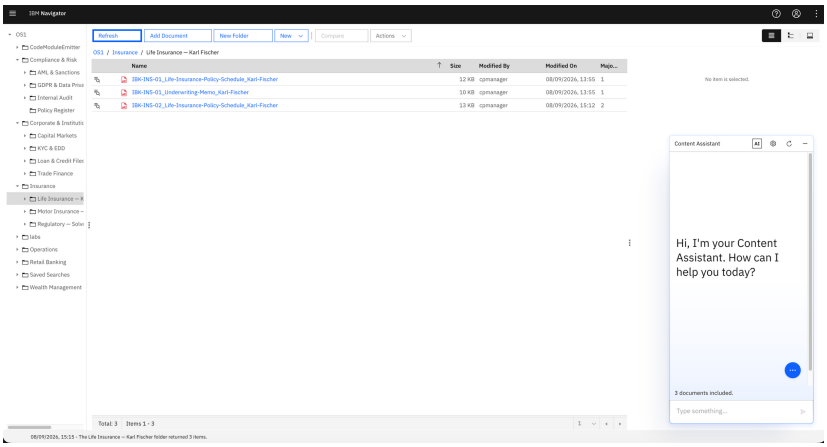

- Select the 2 life insurance policy documents.

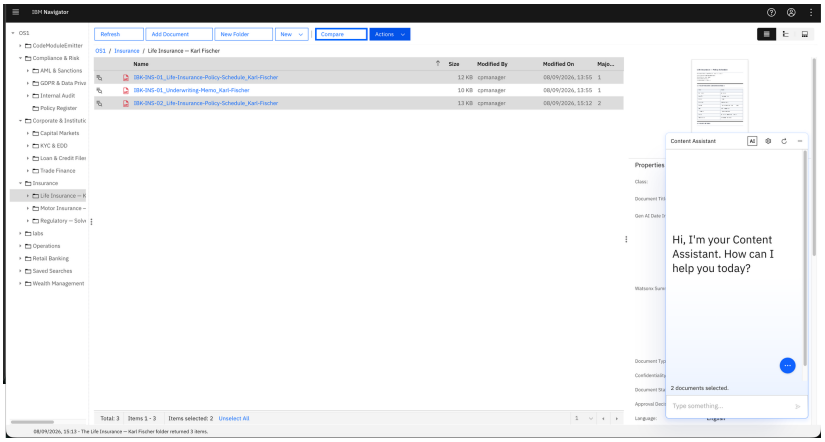

- Click Compare:

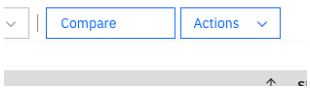

- The assistant can compare selected documents

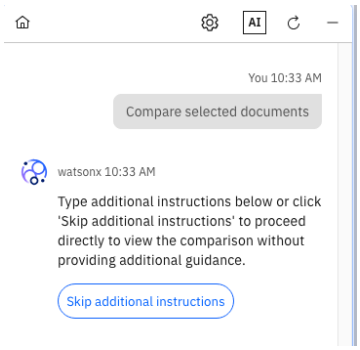

- We can provide further instructions so the Content Assistant can adapt the answer according to our needs. For now, we will skip this by clicking the 'skip additional instructions' button.
- The Content Assistant will summarise the differences.

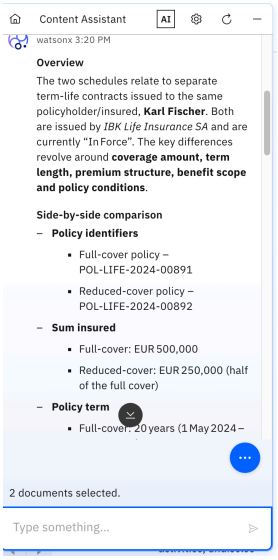

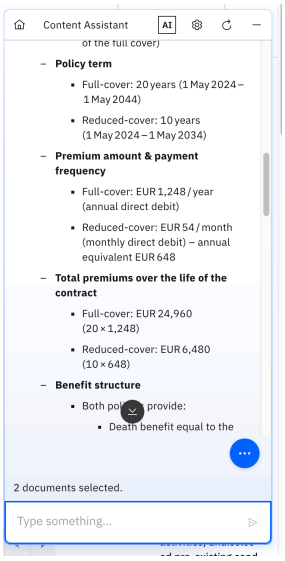

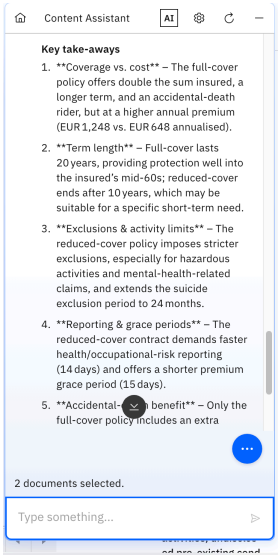

.

IBM® Content Assistant provides natural language insights about business content in your IBM Content Cortex repository. The Content Assistant helps you to quickly find information across large volumes of content and saves time by avoiding manual, repetitive review of document content.

IBM Content Assistant includes generative AI capabilities that can summarise document content, compare documents and versions, and provide accurate answers to questions within the context of the selected business content in an IBM Content Cortex repository.

OK, let’s move to the next level.

4 Lab Exercise 2: Navigator Core Agents

4.1 Introduction

With our built-in agent, you can use natural language to interact with content and perform content management tasks without using traditional interfaces. The agent can also help you organise documents into workspaces—focused collections that provide a dedicated environment for reviewing, analysing, and managing documents identified through your conversations.

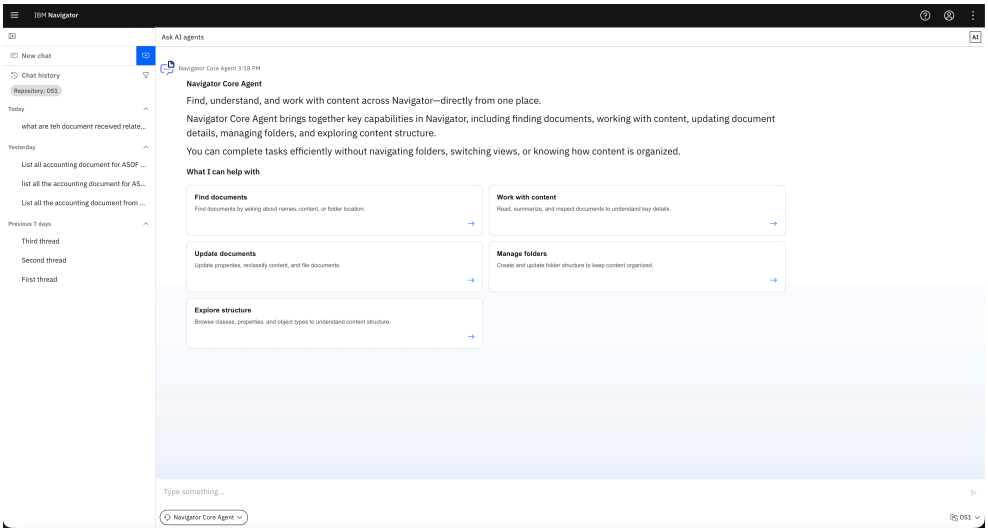

The Navigator Core Agent supports the following content management operations:

- Search for documents by name, content, or folder location
- Create, manage, and navigate folder structures
- View and update document properties and metadata
- Classify documents and get classification recommendations
- Build workspaces to organize documents for review and analysis
- Explore available document classes and their properties

Use the Navigator Core Agent when you need to search for specific documents, organize them into folders, update properties, or build workspaces for focused review and analysis.

Chat history is persisted in the object store and is specific to each object store. When you select an object store, its chat history loads, allowing you to resume previous conversations. Each object store maintains its own separate chat history, and conversations from one object store are not visible when you work with a different object store.

4.2 Technical overview

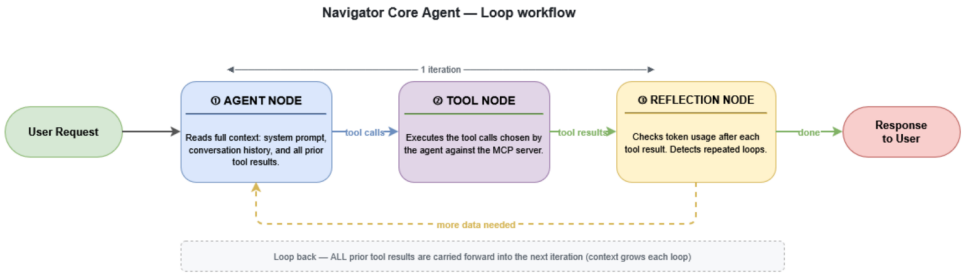

1. **User Request**: The process starts when a user submits a query.
2. ① **Agent Node**:
   - Reads the system prompt, conversation history, and any prior tool results.
   - Decides which tools to call.
3. ② **Tool Node**:
   - Executes the requested tool calls against the **MCP (Model Context Protocol) server** and returns the outputs.
4. ③ **Reflection Node**:
   - Evaluates the tool results.
   - Monitors token usage and checks for infinite / repetitive loops.
   - **Branching Decision**:
     - **More data needed**: Loops back to the **Agent Node**, carrying all prior tool outputs into context.
     - **Done**: Exits the loop and generates the **Response to User**.

4.3 Tour of the new features

4.3.1 Let's switch to Ask AI Agents:

- Navigate to the new Ask AI Agents feature.

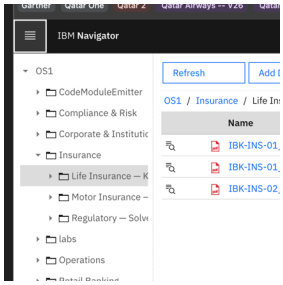

- Click the burger menu, then select Ask AI Agents.

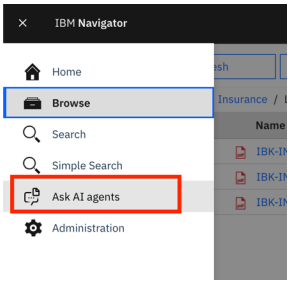

4.4 IBM Navigator: Ask AI Agents

Ask AI Agents is an intelligent conversational framework built on top of your content repository. Instead of navigating folders and running manual searches, you can type natural-language requests, and the agent will query, retrieve, and present your documents instantly. The framework is extensible: you can connect multiple specialised agents to one or more repositories, making it easy to adapt to different business domains and use cases.

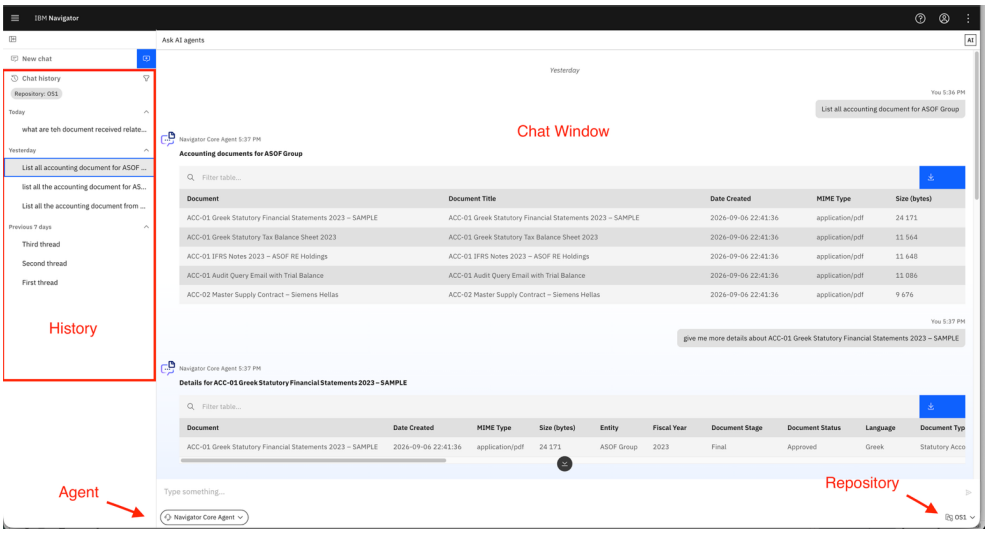

Chat Window (center)

This is the main interaction area where you type your requests and receive responses from the AI agent. The agent displays results directly in the chat, including structured tables with document metadata that you can filter and download. Your messages appear on the right, and the agent's replies appear on the left.

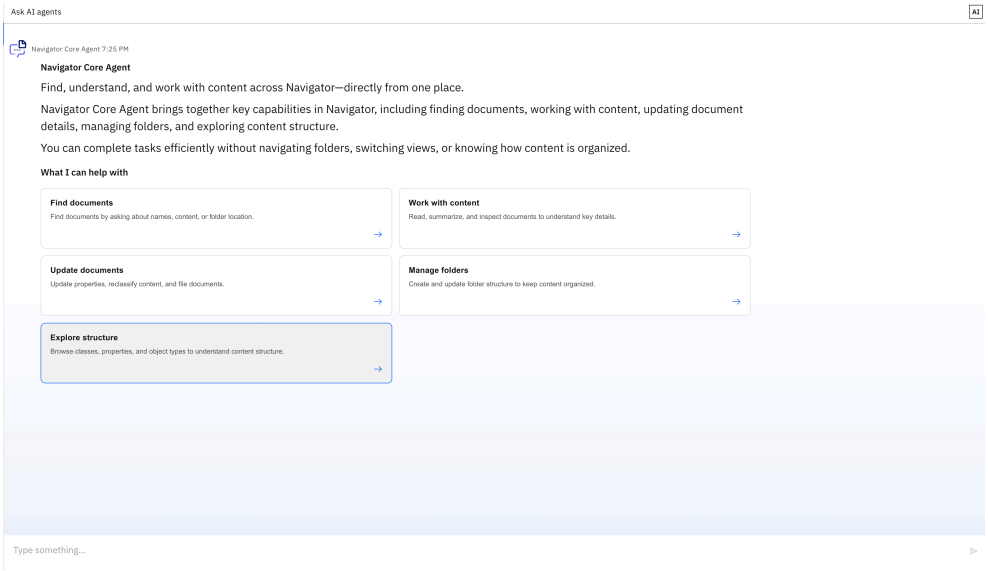

History (left panel)

This panel displays all your past conversations, organized by time period (Today, Yesterday, Previous 7 days). Click any previous thread to revisit it. It also shows the currently connected repository at the top (e.g., *Repository: OS1*).

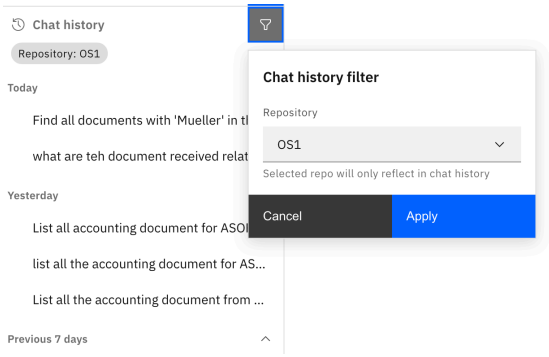

Agent (bottom-left selector)

This selector lets you choose which AI agent handles your requests. In this example, the *Navigator Core Agent* is active. Different agents may have different capabilities.

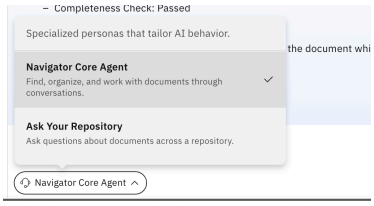

Repository (bottom-right selector)

This indicator shows which content repository the agent currently targets. Here it is set to *OS1*. Changing this value would redirect the agent's queries to a different FileNet object store or repository.

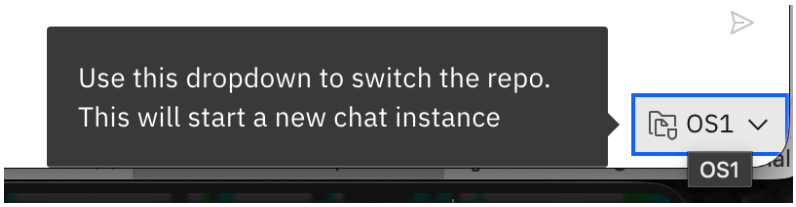

4.4.1 Welcome Screen

When you start a new conversation, the **Navigator Core Agent** greets you with a summary of everything it can do. Rather than requiring you to know where your content lives or how your repository is structured, the agent presents five ready-to-use capability areas—Find documents, Work with content, Update documents, Manage folders, and Explore structure—each accessible by typing a natural language request.

This welcome screen acts as a quick reference, helping both new and experienced users understand the full scope of what the agent can handle before they type their first message.

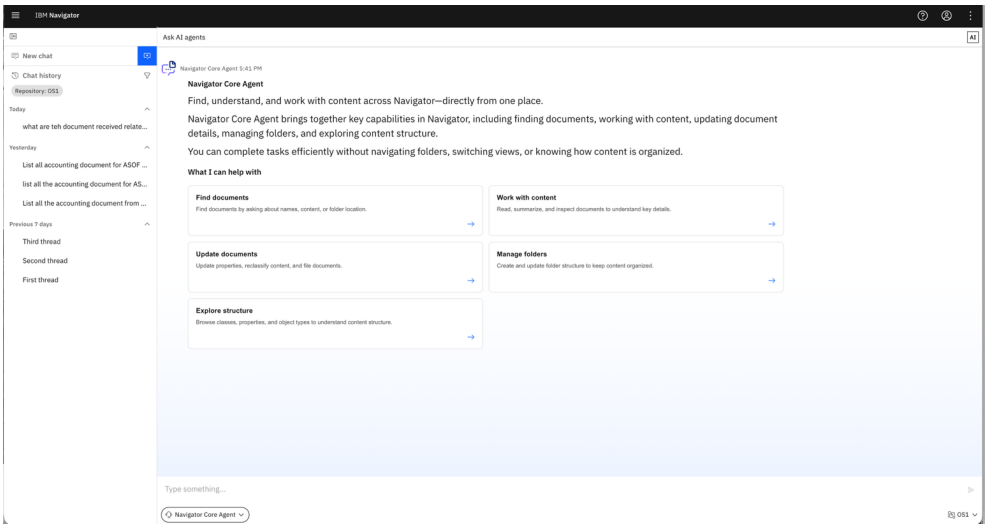

- Click Find documents

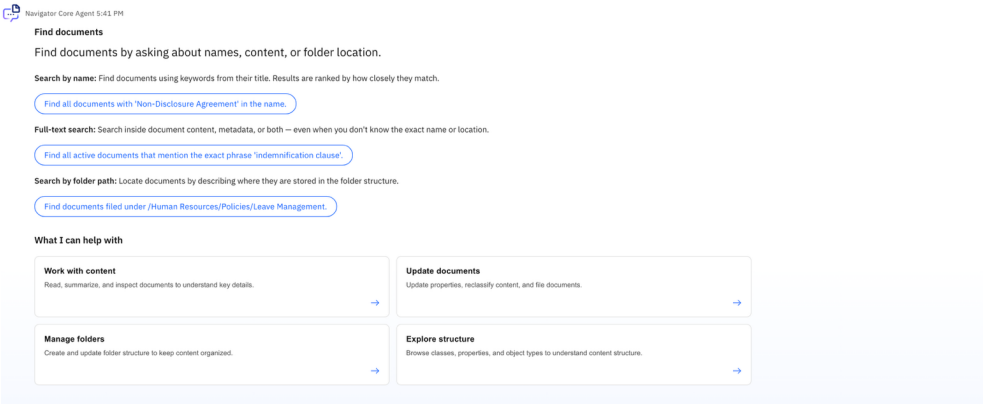

Clicking a capability card expands it into a detailed view, giving you a clear look at how to use that feature with helpful example prompts. The Find Documents capability provides three easy search methods: search by name to quickly find documents using keywords from their titles, full-text search to explore inside document content and metadata even if you're unsure of the exact file name, and search by folder path to find documents based on their storage location within your repository. The clickable example prompts at the bottom of the screen make it simple to get started right away—just click one to use it directly, or customise it to fit your needs.

4.4.2 Find documents

For example, if you need to find all loan documents that mention Mueller. We need to understand the current status and have a small summary for each.

- Ask: “Find all documents with 'Mueller' in the name.”
- You should see the following result:

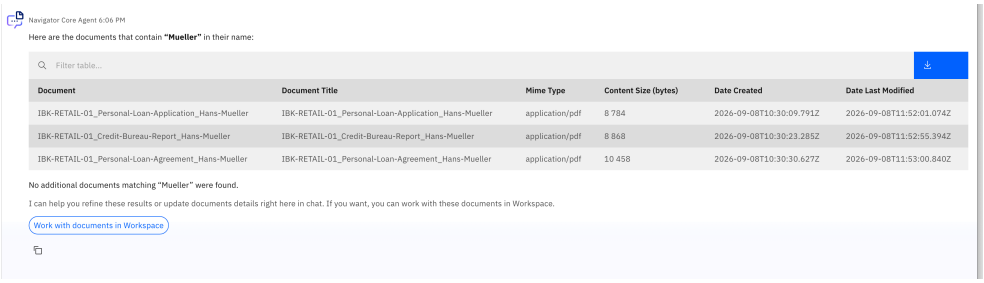

- Ask: “What is the level of confidentiality of each document?

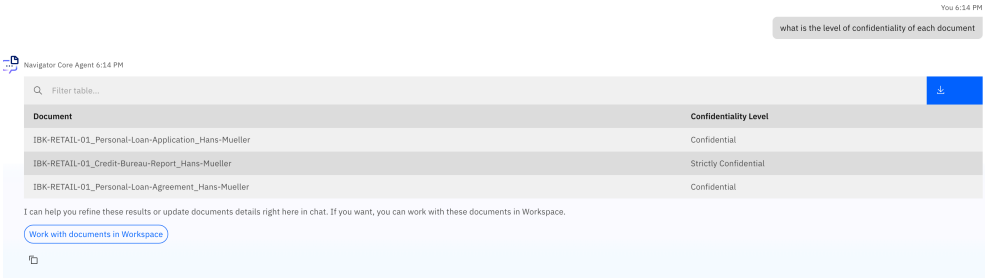

You can update the search results with metadata as needed.

- Ask: “Add document type and Document Status”

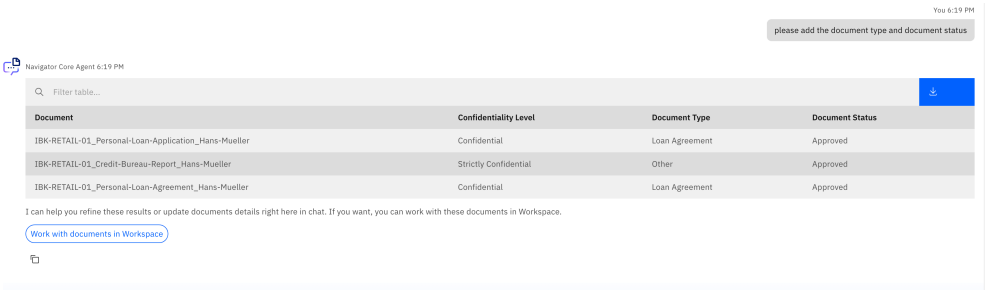

- You can even augment search results with information that is inside the document
- Ask : “Include the loan request and Total Amount Repayable”

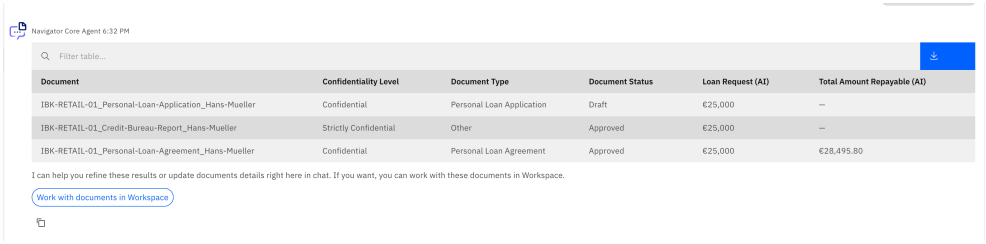

You will notice that next to Loan Request and Total Amount Repayable, an (AI) has been added. This highlights that this information was extracted using Generative AI.

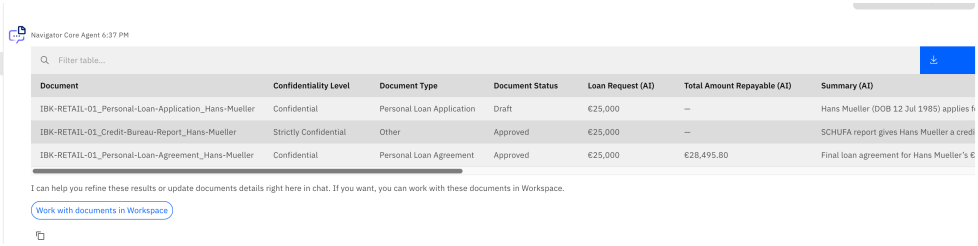

- At any time, you can download the extraction results.

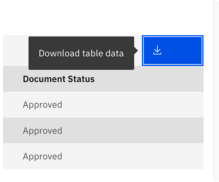

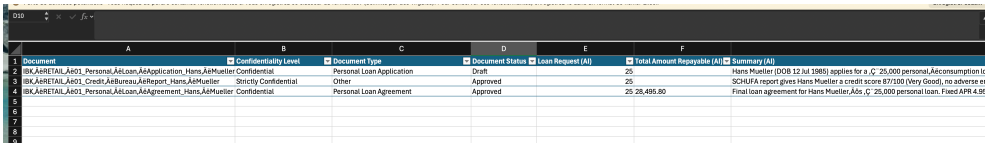

- You may want to access the documents from the search results and be able to perform actions on them. This is what Work with documents in Workspace is all about.

The **Workspace** view opens alongside the chat panel and displays search results in a rich, interactive table. Beyond standard metadata such as document type, status, and confidentiality level, the agent enriches each result with AI-generated columns — including **Loan Request** and **Total Amount Repayable** extracted directly from document content, and a **Summary (AI)** column that provides a concise plain-language synopsis of each document. This lets users grasp the key facts of multiple documents at a glance, without opening a single file. From the Workspace, you can also continue refining results or take actions on selected documents directly from the chat.

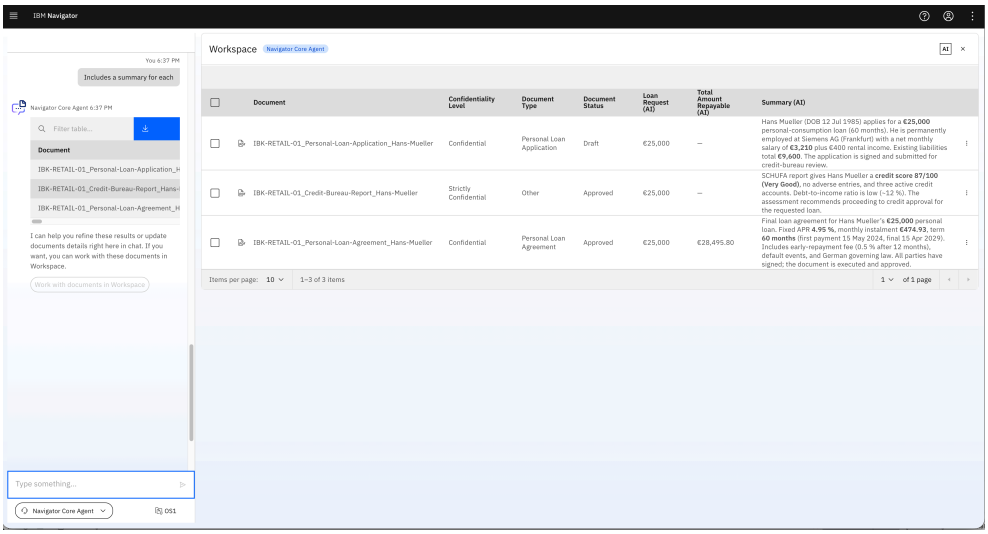

- You can perform an action on a single document by selecting the action menu in that document's row ⋮

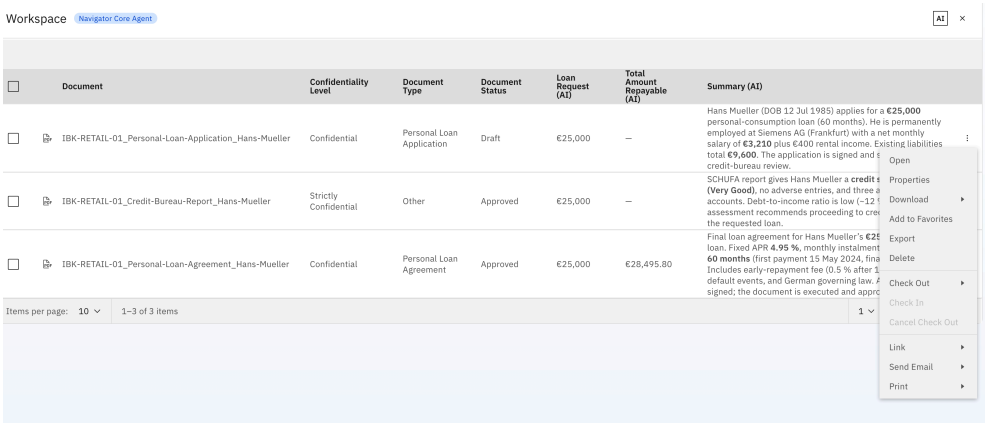

- The Workspace also supports bulk actions. When you select multiple documents using the checkboxes, a blue action bar appears at the top indicating how many items are selected, and exposes options to **Export** or perform additional actions via the **Menu**. This makes it straightforward to package a set of related documents — for example, all files belonging to a loan application — and export them in one operation, without ever leaving the chat interface.

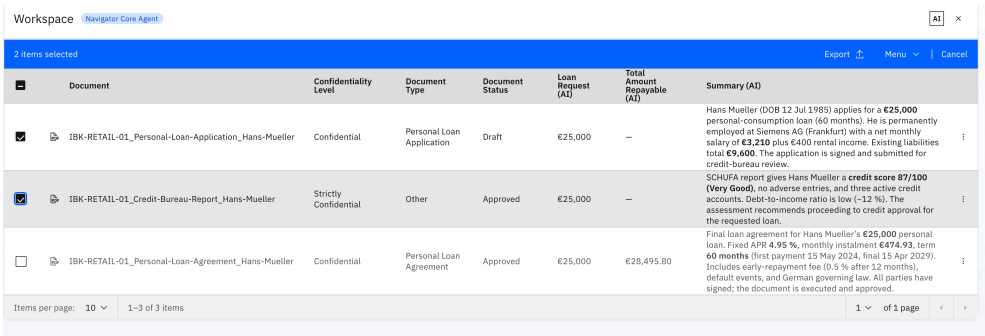

- Select the first document on the row; you will notice that the agent has that action in its context.

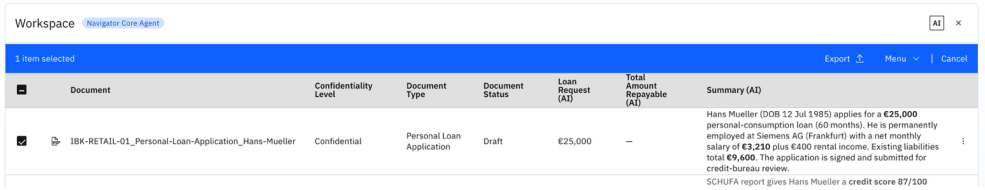

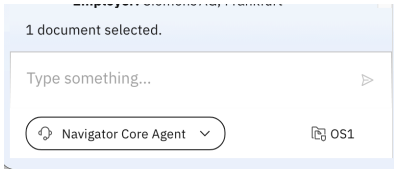

- Ask: “What is the key information?”

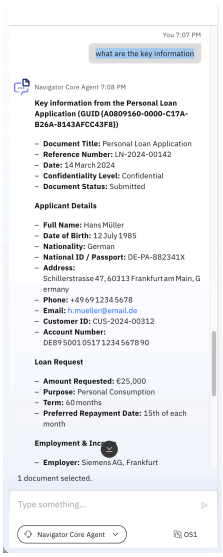

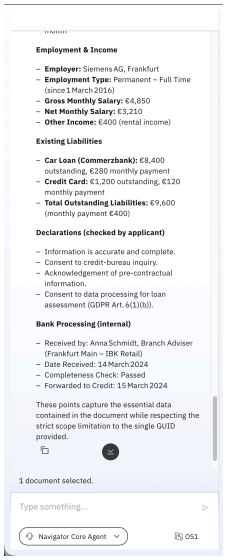

- Close the workspace.

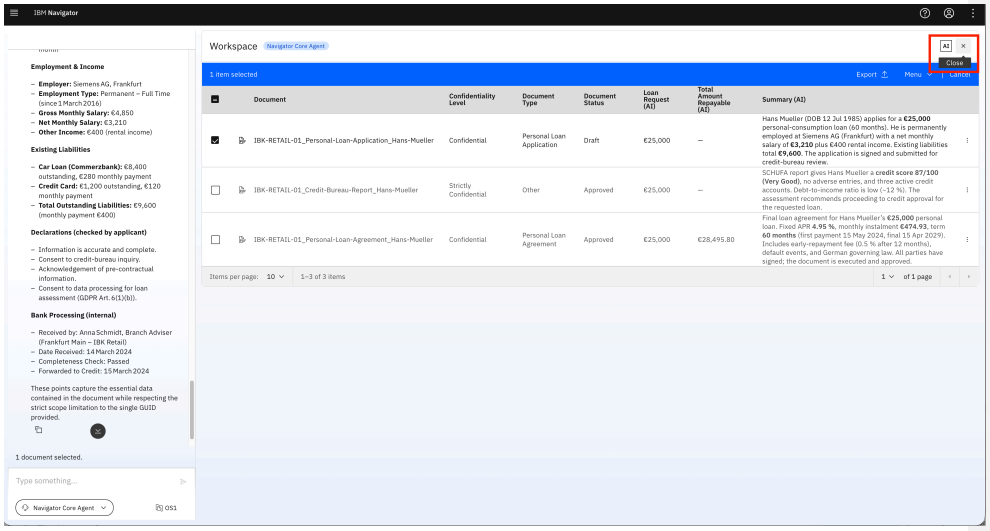

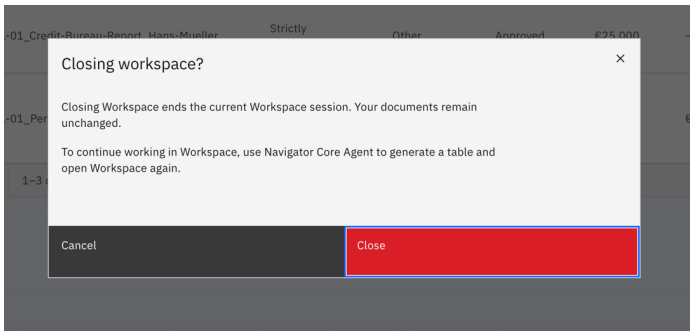

4.4.3 Update documents

We received a set of documents, and we will use our Navigator Core Agent to help us classify and organise them. You will be working in parallel on the same system. You will receive a reference ID (User001-User020) from the Lab Instructor. First, we will try to understand the current document class definition and map it to the document received.

- Start browse feature

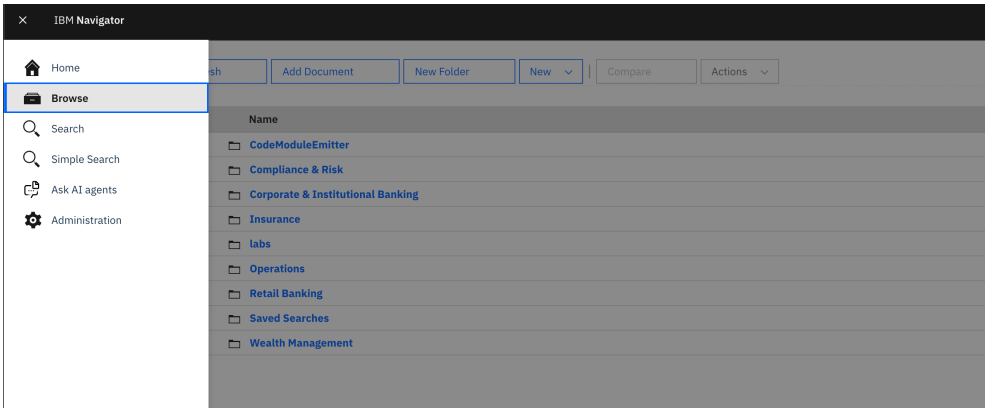

- Navigate to your user lab folder. You will notice that all the documents have been classified as the Document class. It's a generic type.

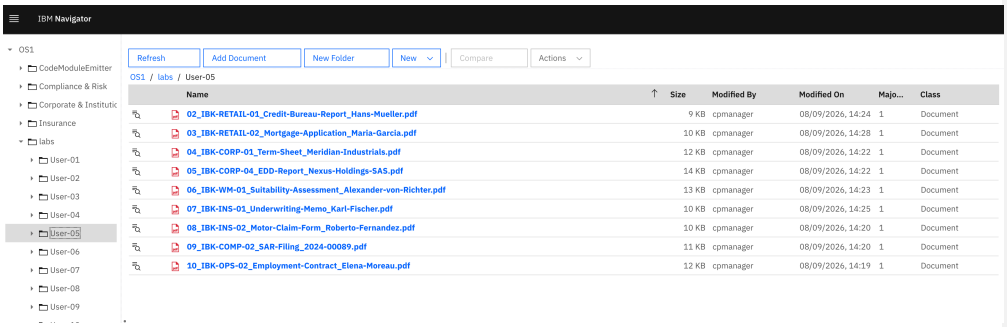

- Let's start a new chat

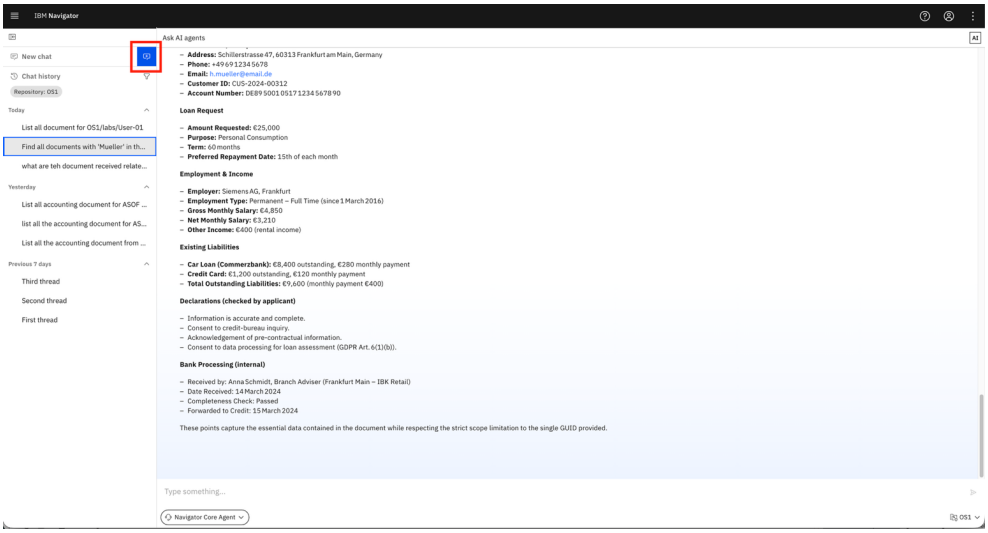

- A new welcome screen should be available:

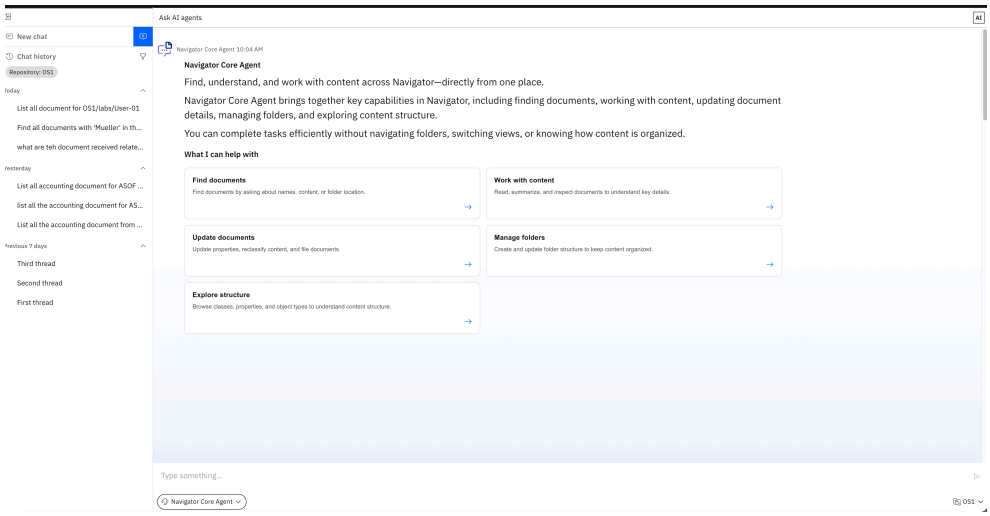

- Click on “Update Documents”

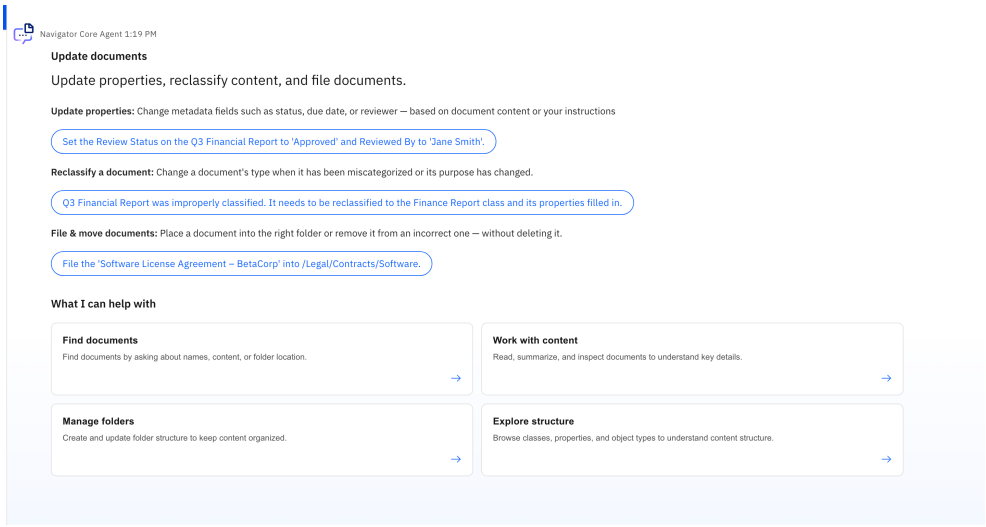

- The AI agent proposes a couple of sample commands that support the update Document scenario.
- Ask: “+”.

Navigator Core Agent 8:38 PM

IBK-related Document Types

Filter table...

Class Display Name	Symbolic Name	Description
IBK Compliance Document	IBKComplianceDocument	Compliance & Risk document class covering regulatory policy, sanctions screening, internal audit, and GDPR.
IBK Corporate Document	IBKCorporateDocument	Corporate & Institutional Banking document class covering syndicated loans, trade finance, capital markets, and KYC/EDD.
IBK Document	IBKDocument	Base document class for all IBK Insurance & Banking line-of-business document subclasses.
IBK Insurance Document	IBKInsuranceDocument	Insurance document class covering policy lifecycle, claims, reinsurance, and regulatory reporting.
IBK Operations Document	IBKOperationsDocument	Operations & Back-Office document class covering vendor contracts, invoices, HR records, and legacy migration.
IBK Retail Document	IBKRetailDocument	Retail Banking document class covering personal loans, mortgages, KYC/AML, and account statements.
IBK Wealth Document	IBKWealthDocument	Wealth Management document class covering private client onboarding, portfolio mandates, and trust administration.

- Ask: “List all the documents for /labs/User-0x”, The X need to be replaced based on the ID provided by the instructor

You 8:36 PM

List all document for OS1/labs/User-01

Navigator Core Agent 8:36 PM

Documents in /labs/User-01 (OS1 repository)

Filter table...

Document
02_IBK-RETAIL-01_Credit-Bureau-Report_Hans-Mueller.pdf
03_IBK-RETAIL-02_Mortgage-Application_Maria-Garcia.pdf
04_IBK-CORP-01_Term-Sheet_Meridian-Industrials.pdf
05_IBK-CORP-04_EDD-Report_Nexus-Holdings-SAS.pdf
06_IBK-WM-01_Suitability-Assessment_Alexander-von-Richter.pdf
07_IBK-INS-01_Underwriting-Memo_Karl-Fischer.pdf
08_IBK-INS-02_Motor-Claim-Form_Roberto-Fernandez.pdf
09_IBK-COMP-02_SAR-Filing_2024-00089.pdf
10_IBK-OPS-02_Employment-Contract_Elena-Moreau.pdf

- Ask: “Provide me a resume for the first document and tell me it's classified correctly”

Navigator Core Agent 10:13 AM

Resume (summary) of the document

- **Title:** 03 IBK-RETAIL-02 Mortgage-Application Maria-Garcia.pdf
- **Type:** Mortgage application (IBK-RETAIL-02) – submitted 5 Apr 2024 to IBK Bank, Madrid branch.
- **Applicant:** María García López (DOB 23 Nov 1982, DNI 12345678A, married to Javier García Ruiz, DNI 87654321B).
- **Property:** Residential apartment, Calle Mayor 15, 1ºA, Madrid; purchase price €380 000, mortgage request €285 000 (75 % LTV), 25-year fixed-rate.
- **Income & Employment:**
  - María – Senior Manager, Deloitte España, €72 000 gross/year, €4 100 net/month.
  - Javier – Self-employed architect, €48 000 annual declared income, €2 900 net/month.
- **Existing commitments:** €470 /month (car loan + credit-card).
- **Documents submitted:** DNI copies, last 3 payslips (main applicant), 2022-23 tax returns (both), 6-month bank statements, preliminary property extract (Nota Simple). Valuation report pending.
- **Status:** Draft (IBK-DocumentStatus = Draft).

Classification check

The document is currently stored with the generic class **Document**.
Based on its content and the IBK-related document-type list, it belongs to the **IBK Retail Document** class (the “IBK-RETAIL-02” subtype).

Conclusion: The document is **not classified correctly**; it should be re-classified to **IBK Retail Document** to reflect its true type.

- Ask: “Reclassify the document into the retail document.”

- Answer: Yes

- Ask : “Update the Document Status to Under Review, Customer Name to Maria Garcia, and change the classification to Confidential”

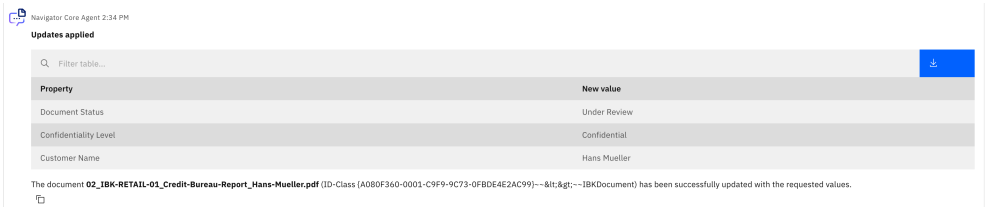

- Answer: OK
- Ask: Search all folder that contains Maria Garcia

- Ask to move the document to that folder.

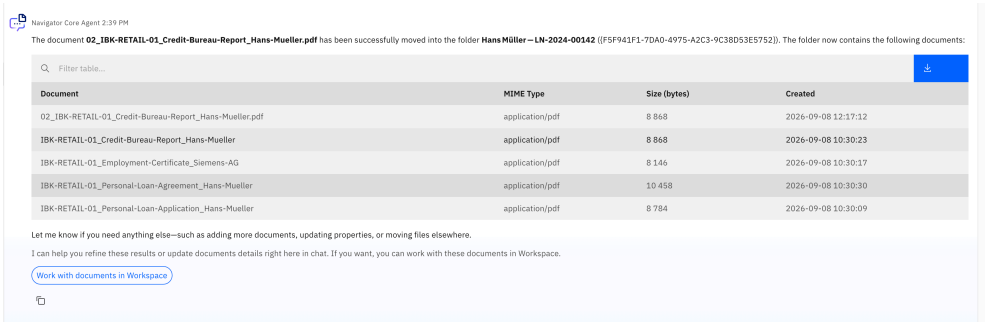

- Congratulations, you classified and moved the document to the right directory.
- Let’s move back to our User-0X folder.

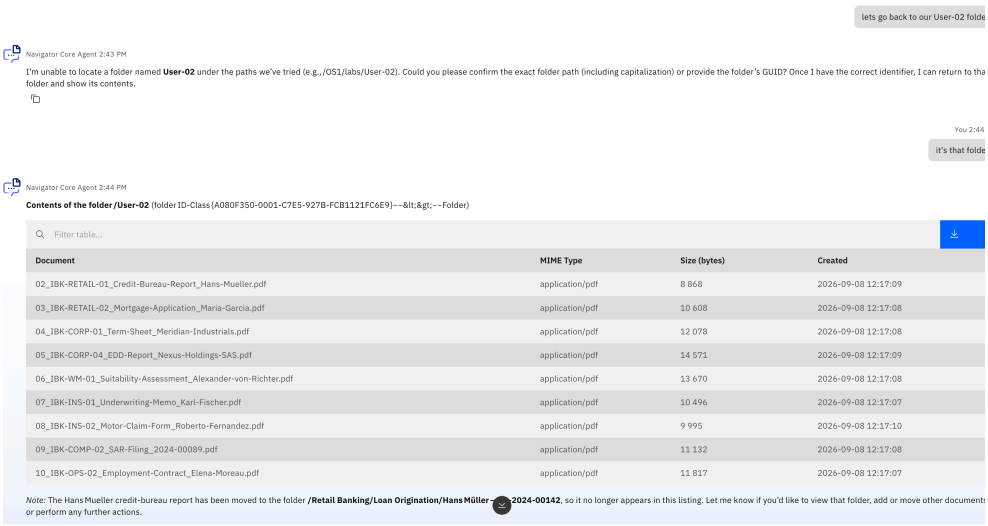

- Instead of classifying documents one by one, let’s move to mass classification.
- Ask: “Can you review all the remaining documents in that folder and recommend the correct classification for each one ?”

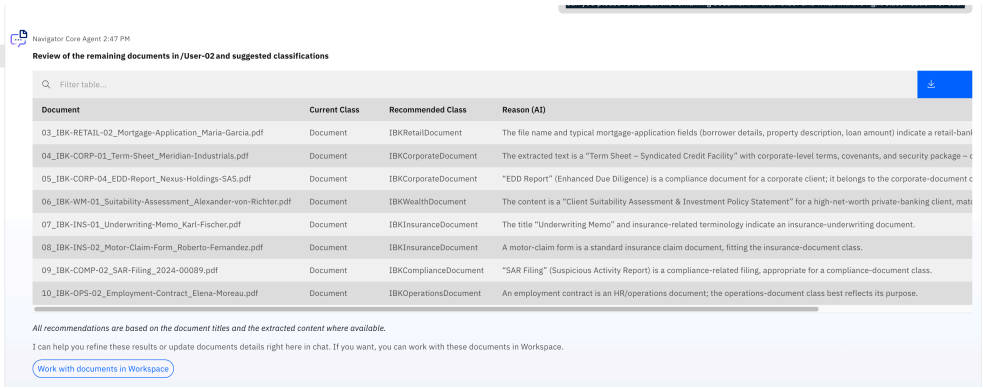

Navigator Core Agent 2:47 PM

Review of the remaining documents in /User-02 and suggested classifications

Filter table...

Document	Current Class	Recommended Class	Reason (AI)
03_IBK-RETAIL-02_Mortgage-Application_Maria-Garcia.pdf	Document	IBKRetailDocument	The file name and typical mortgage-application fields (borrower details, property description, loan amount) indicate a retail-bank
04_IBK-CORP-01_Term-Sheet_Meridian-Industrials.pdf	Document	IBKCorporateDocument	The extracted text is a "Term Sheet – Syndicated Credit Facility" with corporate-level terms, covenants, and security package – c
05_IBK-CORP-04_EDD-Report_Nexus-Holdings-SAS.pdf	Document	IBKCorporateDocument	"EDD Report" (Enhanced Due Diligence) is a compliance document for a corporate client; it belongs to the corporate-document c
06_IBK-WM-01_Suitability-Assessment_Alexander-von-Richter.pdf	Document	IBKWealthDocument	The content is a "Client Suitability Assessment & Investment Policy Statement" for a high-net-worth private-banking client, mati
07_IBK-INS-01_Underwriting-Memo_Karl-Fischer.pdf	Document	IBKInsuranceDocument	The title "Underwriting Memo" and insurance-related terminology indicate an insurance-underwriting document.
08_IBK-INS-02_Motor-Claim-Form_Roberto-Fernandez.pdf	Document	IBKInsuranceDocument	A motor-claim form is a standard insurance claim document, fitting the insurance-document class.
09_IBK-COMP-02_SAR-Filing_2024-00089.pdf	Document	IBKComplianceDocument	"SAR Filing" (Suspicious Activity Report) is a compliance-related filing, appropriate for a compliance-document class.
10_IBK-OPS-02_Employment-Contract_Elena-Moreau.pdf	Document	IBKOperationsDocument	An employment contract is an HR/operations document; the operations-document class best reflects its purpose.

All recommendations are based on the document titles and the extracted content where available.

I can help you refine these results or update documents details right here in chat. If you want, you can work with these documents in Workspace.

Work with documents in Workspace

- We can review the list of documents, the proposed document class for each, and the reason for the recommendation.
- Answer: "OK"

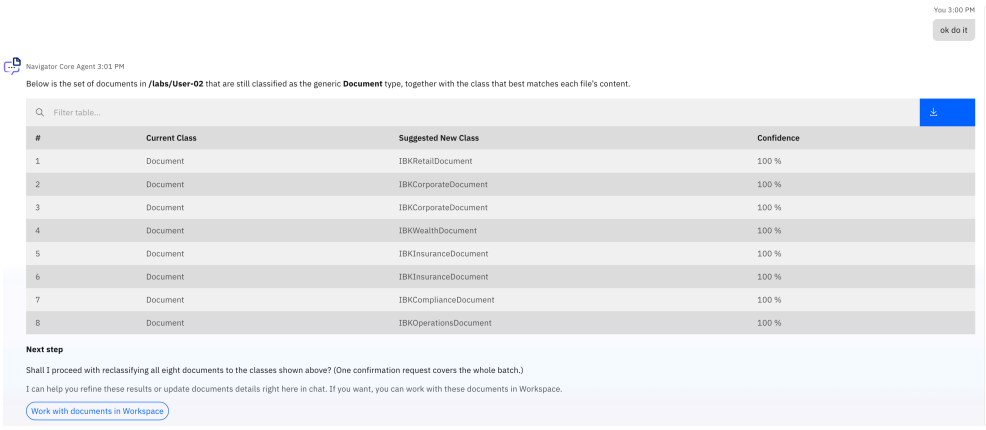

You 3:00 PM

ok do it

Navigator Core Agent 3:01 PM

Below is the set of documents in **/labs/User-02** that are still classified as the generic **Document** type, together with the class that best matches each file's content.

Filter table...

#	Current Class	Suggested New Class	Confidence
1	Document	IBKRetailDocument	100 %
2	Document	IBKCorporateDocument	100 %
3	Document	IBKCorporateDocument	100 %
4	Document	IBKWealthDocument	100 %
5	Document	IBKInsuranceDocument	100 %
6	Document	IBKInsuranceDocument	100 %
7	Document	IBKComplianceDocument	100 %
8	Document	IBKOperationsDocument	100 %

Next step

Shall I proceed with reclassifying all eight documents to the classes shown above? (One confirmation request covers the whole batch.)

I can help you refine these results or update documents details right here in chat. If you want, you can work with these documents in Workspace.

Work with documents in Workspace

- OK, reclassify them.

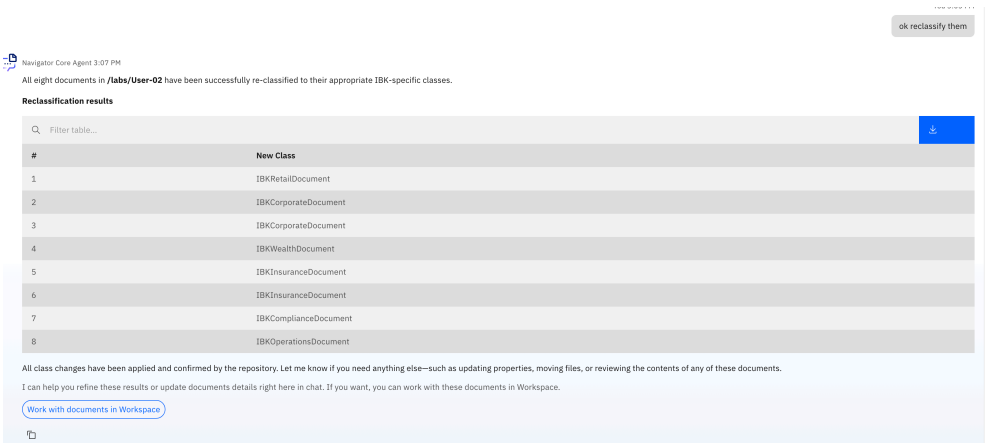

- Let’s switch to our Browse feature to review the changes in the folder:

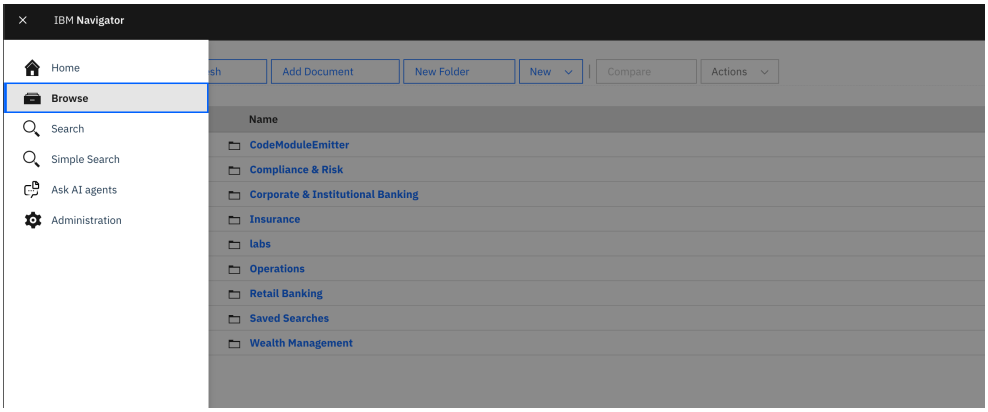

- Navigate to our lab user to notice that the class has been updated to the correct class

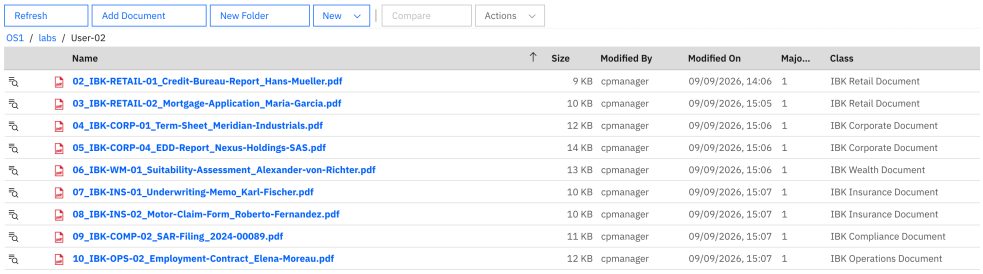

Congratulations—you have updated eight documents in a single action.

In the upcoming lab, we will utilize a third-party agent to update the metadata and organize them into the correct directory.

5 Lab Exercise 3: Integrate Content Cortex capabilities

5.1 Content Cortex Agentic Integration:

The **Model Context Protocol (MCP)** capability in Content Cortex gives third-party AI agents — built on any framework — a standardised way to interact with IBM's content management platform. No bespoke integrations required.

What Agents Can Do

Connected agents can search and retrieve documents, create and version new content, manage folder structures, and extract document text to feed directly into LLM prompts. Every operation flows through the same Content Engine security model, so governance and audit trails remain intact.

AI-Driven Classification

The most immediate unlock is **automated classification**. An agent reads a document's text via MCP, sends it to a language model, then writes the inferred class and metadata properties back — closing the loop between ingestion and governance without human intervention.

The Result

Content Cortex MCP turns the repository from a passive file store into an active, queryable knowledge layer that AI agents can read from, write to, and govern against — through a single, open protocol.

5.1.1 Let's get the lab starter :

- Download the lab zip folder in the box folder:
- Open the project in bob :

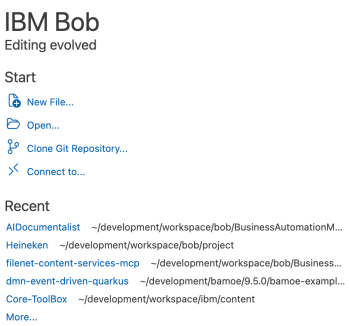

- You will notice that there is a .bob folder that contains an mcp.json configuration

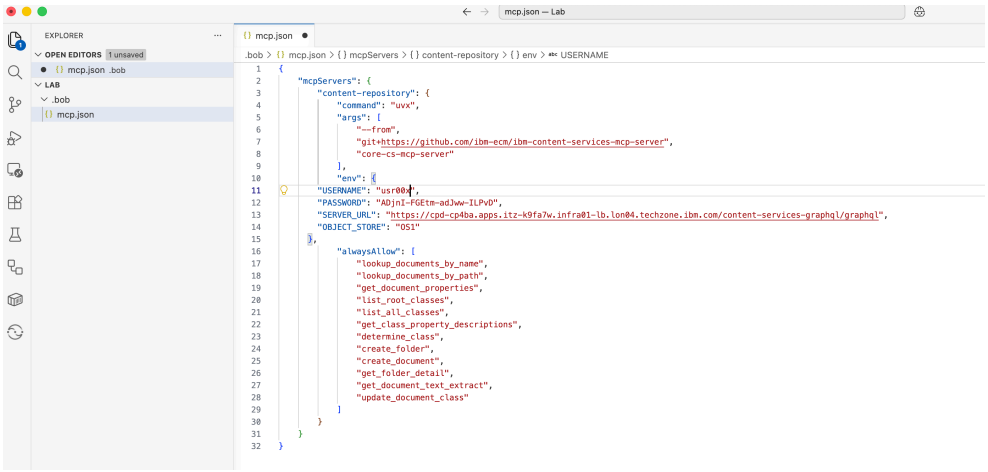

- Change the username to the one assigned to you by the instructor.
- Expand the bob chat window to take full screen.

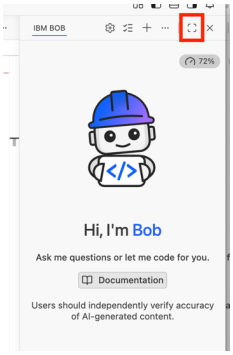

- Ask : “list all the documents under /OS1/labs/User-02”

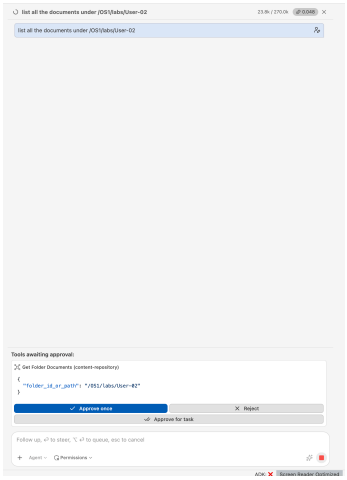

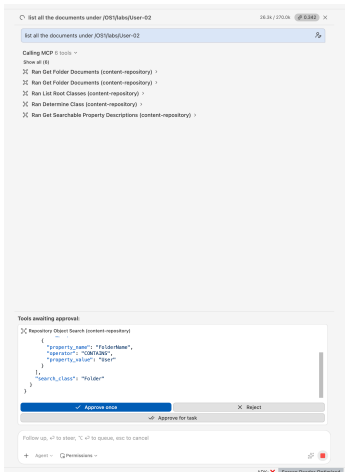

- Accept all the sub-actions that the system asks you to be able to process the query.

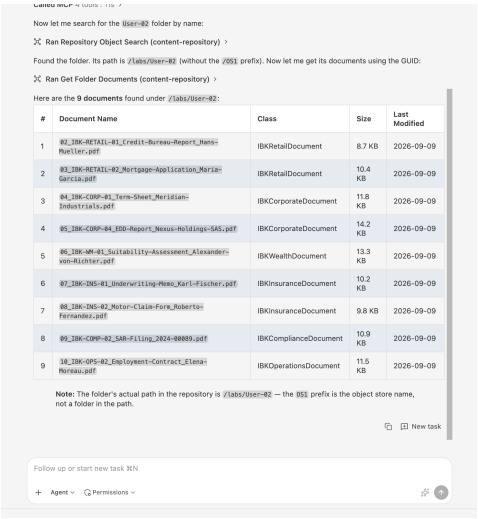

- Ask: “Review the document type datamodel, review document content and complete the metadata around the document for the first document.”

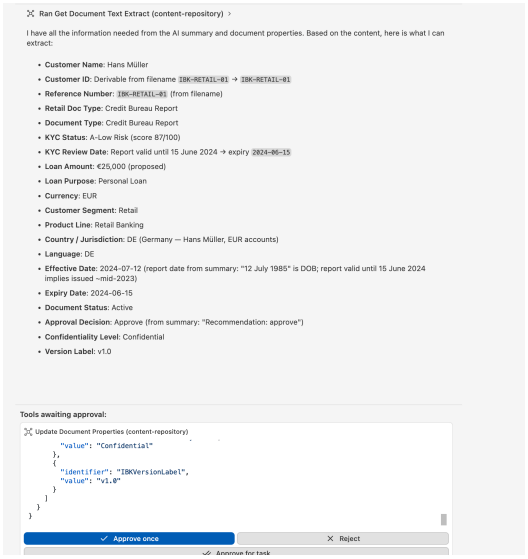

- Approve to update the document definition. And he will notice some constraint for some properties:

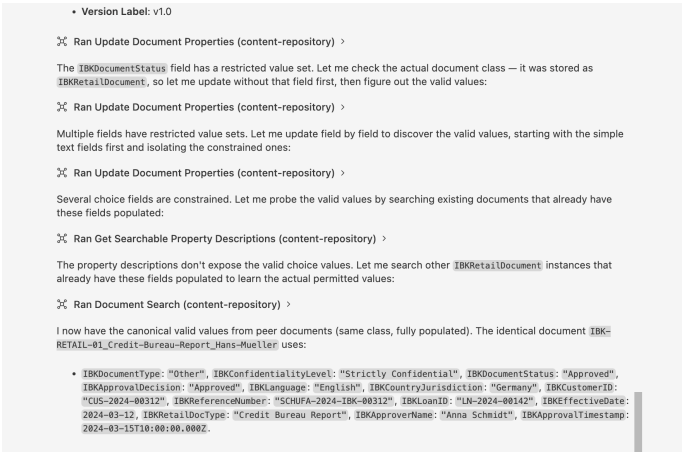

- Document has been updated :

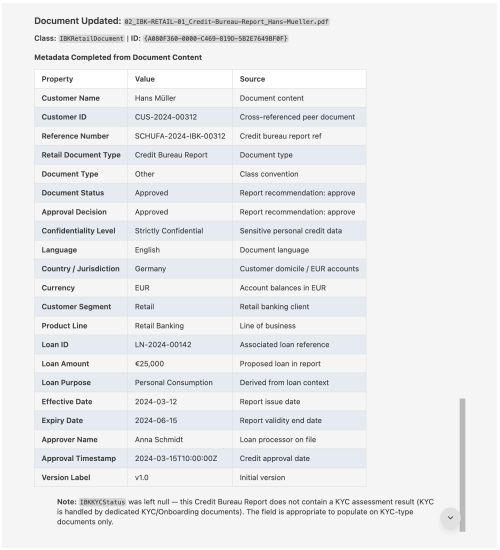

Document Updated: `02_IBK-RETAIL-01_Credit-Bureau-Report_Hans-Mueller.pdf`

Class: `IBKRetailDocument` | **ID:** `{A8B0F360-0000-C469-B19D-5B2E7649BF0F}`

Metadata Completed from Document Content

Property	Value	Source
Customer Name	Hans Müller	Document content
Customer ID	CUS-2024-00312	Cross-referenced peer document
Reference Number	SCHUFA-2024-IBK-00312	Credit bureau report ref
Retail Document Type	Credit Bureau Report	Document type
Document Type	Other	Class convention
Document Status	Approved	Report recommendation: approve
Approval Decision	Approved	Report recommendation: approve
Confidentiality Level	Strictly Confidential	Sensitive personal credit data
Language	English	Document language
Country / Jurisdiction	Germany	Customer domicile / EUR accounts
Currency	EUR	Account balances in EUR
Customer Segment	Retail	Retail banking client
Product Line	Retail Banking	Line of business
Loan ID	LN-2024-00142	Associated loan reference
Loan Amount	€25,000	Proposed loan in report
Loan Purpose	Personal Consumption	Derived from loan context
Effective Date	2024-03-12	Report issue date
Expiry Date	2024-06-15	Report validity end date
Approver Name	Anna Schmidt	Loan processor on file
Approval Timestamp	2024-03-15T10:00:00Z	Credit approval date
Version Label	v1.0	Initial version

Note: `IBKKYCStatus` was left null — this Credit Bureau Report does not contain a KYC assessment result (KYC is handled by dedicated KYC/Onboarding documents). The field is appropriate to populate on KYC-type documents only.

- Ask: “Apply the same logic for the rest of the documents”

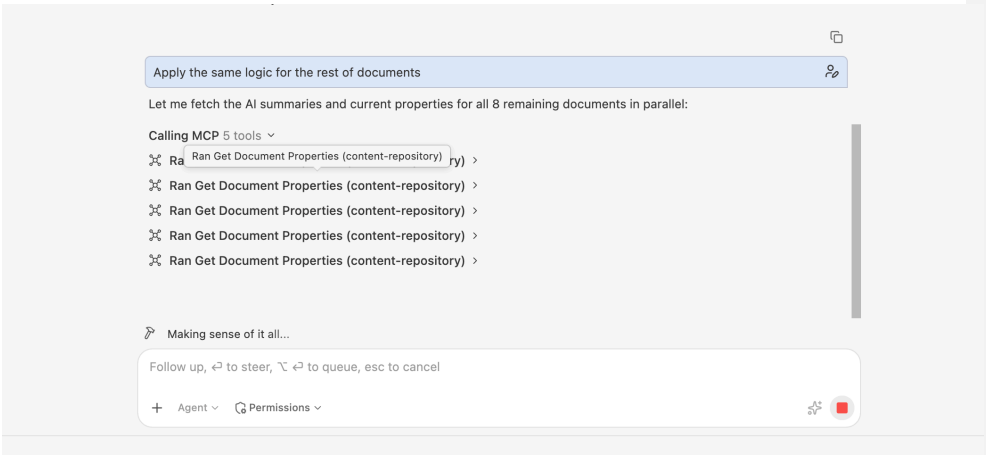

- After a couple of seconds:

All 9 Documents in `/labs/User-02` — Metadata Completed ✅

#	Document	Class	Key Fields Set
1	`02_IBK-RETAIL-01_Credit-Bureau-Report_Hans-Mueller.pdf`	IBKRetailDocument	CustomerName, CustomerID, RefNo, DocType, KYCStatus, LoanAmount, Currency, Jurisdiction, Status=Approved *(done previously)*
2	`03_IBK-RETAIL-02_Mortgage-Application_Maria-Garcia.pdf`	IBKRetailDocument	CustomerName=María García López, LoanAmount=€285k, CollateralType=Residential Property, Status=Under Review, Jurisdiction=Spain
3	`04_IBK-CORP-01_Term-Sheet_Meridian-Industrials.pdf`	IBKCorporateDocument	CustomerName=Meridian Industrials GmbH, FacilityAmount=€150M, FacilityType=Syndicated Loan, MarginBPS=175, EDDRiskRating=Medium
4	`05_IBK-CORP-04_EDD-Report_Nexus-Holdings-SAS.pdf`	IBKCorporateDocument	CustomerName=Nexus Holdings SAS, DocType=KYC/AML Record, EDDRiskRating=**High**, CorpDocType=EDD Report
5	`06_IBK-WM-01_Suitability-Assessment_Alexander-von-Richter.pdf`	IBKWealthDocument	CustomerName=Alexander von Richter, Segment=High Net Worth, RiskProfile=Growth, AUMValue=€5.2M, PortfolioID=PF-2024-0031
6	`07_IBK-INS-01_Underwriting-Memo_Karl-Fischer.pdf`	IBKInsuranceDocument	InsuredName=Karl Fischer, PolicyNo=POL-LIFE-2024-00891, Product=Term Life, SumInsured=€500k, Decision=Accepted Standard
7	`08_IBK-INS-02_Motor-Claim-Form_Roberto-Fernandez.pdf`	IBKInsuranceDocument	InsuredName=Roberto Fernández Navarro, ClaimID=CLM-2024-04421, ClaimAmount=€4,800, ClaimStatus=Registered, LossDate=2024-05-11
8	`09_IBK-COMP-02_SAR-Filing_2024-00089.pdf`	IBKComplianceDocument	CaseID=SAR-2024-IBK-00089, CompDocType=SAR Filing, ComplianceStatus=Regulatory Submission, Regulator=FIU Germany, Confidentiality=Strictly Confidential
9	`10_IBK-OPS-02_Employment-Contract_Elena-Moreau.pdf`	IBKOperationsDocument	EmployeeName=Elena Moreau, EmployeeID=EMP-2024-0817, OpsDocType=Employment Contract, Department=Digital Product Management, ContractStatus=Active

New task

- The second part of our journey is to move the document to the right folder.
- Ask: ” Based on the IBK management plan, where the first document needs to be moved“

Document Moved per IBK ECM Management Plan

Document: `02_IBK-RETAIL-01_Credit-Bureau-Report_Hans-Mueller.pdf`

Plan reference: Sub-task 2 — Retail Banking, folder hierarchy rule:

> Credit Bureau Reports → `/Retail Banking/Loan Origination/{Customer Name} – {Loan ID}/03 – Credit Assessment/`

Previous location	`/labs/User-02`
Target folder	`/Retail Banking/Loan Origination/Hans Müller – LN-2024-00142/03 – Credit Assessment/`
Action	Sub-folder `03 – Credit Assessment` created (didn't exist yet), document filed into it

> Note: The document remains filed in `/labs/User-02` as well (content repository uses referential containment — filing doesn't move, it adds a reference). If you want it removed from `/labs/User-02`, I can unfile it from there.

New task

- Perfect, next step: “Do it for the rest of the documents”

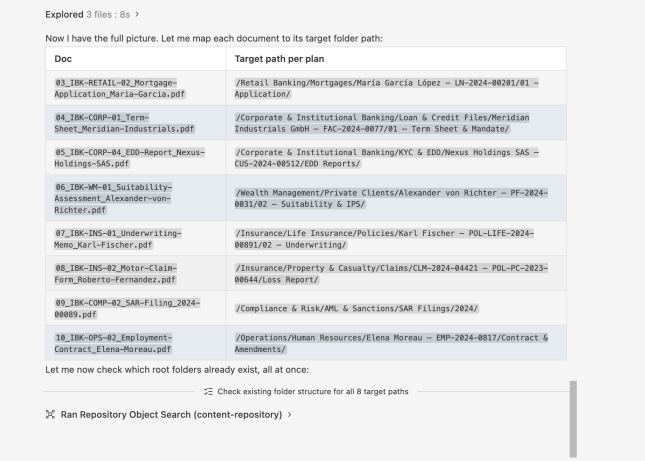

Explored 3 files · 8s ›

Now I have the full picture. Let me map each document to its target folder path:

Doc	Target path per plan
`03_IBK-RETAIL-02_Mortgage-Application_Maria-Garcia.pdf`	`/Retail Banking/Mortgages/María García López – LN-2024-00201/01 – Application/`
`04_IBK-CORP-01_Term-Sheet_Meridian-Industrials.pdf`	`/Corporate & Institutional Banking/Loan & Credit Files/Meridian Industrials GmbH – FAC-2024-0077/01 – Term Sheet & Mandate/`
`05_IBK-CORP-04_EDD-Report_Nexus-Holdings-SAS.pdf`	`/Corporate & Institutional Banking/KYC & EDD/Nexus Holdings SAS – CUS-2024-00512/EDD Reports/`
`06_IBK-WM-01_Suitability-Assessment_Alexander-von-Richter.pdf`	`/Wealth Management/Private Clients/Alexander von Richter – PF-2024-0031/02 – Suitability & IPS/`
`07_IBK-INS-01_Underwriting-Memo_Karl-Fischer.pdf`	`/Insurance/Life Insurance/Policies/Karl Fischer – POL-LIFE-2024-00891/02 – Underwriting/`
`08_IBK-INS-02_Motor-Claim-Form_Roberto-Fernandez.pdf`	`/Insurance/Property & Casualty/Claims/CLM-2024-04421 – POL-PC-2023-00644/Loss Report/`
`09_IBK-COMP-02_SAR-Filing_2024-00089.pdf`	`/Compliance & Risk/AML & Sanctions/SAR Filings/2024/`
`10_IBK-OPS-02_Employment-Contract_Elena-Moreau.pdf`	`/Operations/Human Resources/Elena Moreau – EMP-2024-0817/Contract & Amendments/`

Let me now check which root folders already exist, all at once:

Check existing folder structure for all 8 target paths

Ran Repository Object Search (content-repository) ›

- After a couple of seconds:

All Documents Filed per IBK ECM Management Plan ✅

#	Document	Target Path (created & filed)
2	`03_IBK-RETAIL-02_Mortgage-Application_Maria-Garcia.pdf`	`/Retail Banking/Mortgages/María García López — LN-2024-00201/01 — Application/`
3	`04_IBK-CORP-01_Term-Sheet_Meridian-Industrials.pdf`	`/Corporate & Institutional Banking/Loan & Credit Files/Meridian Industrials GmbH — FAC-2024-0077/01 — Term Sheet & Mandate/`
4	`05_IBK-CORP-04_EDD-Report_Nexus-Holdings-SAS.pdf`	`/Corporate & Institutional Banking/KYC & EDD/Nexus Holdings SAS — CUS-2024-00512/EDD Reports/`
5	`06_IBK-WM-01_Suitability-Assessment_Alexander-von-Richter.pdf`	`/Wealth Management/Private Clients/Alexander von Richter — PF-2024-0031/02 — Suitability & IPS/`
6	`07_IBK-INS-01_Underwriting-Memo_Karl-Fischer.pdf`	`/Insurance/Life Insurance/Policies/Karl Fischer — POL-LIFE-2024-00891/02 — Underwriting/`
7	`08_IBK-INS-02_Motor-Claim-Form_Roberto-Fernandez.pdf`	`/Insurance/Property & Casualty/Claims/CLM-2024-04421 — POL-PC-2023-00644/Loss Report/`
8	`09_IBK-COMP-02_SAR-Filing_2024-00089.pdf`	`/Compliance & Risk/AML & Sanctions/SAR Filings/2024/`
9	`10_IBK-OPS-02_Employment-Contract_Elena-Moreau.pdf`	`/Operations/Human Resources/Elena Moreau — EMP-2024-0817/Contract & Amendments/`

Folders created: 18 new folders across 5 division trees. All documents retain their `/labs/User-02` filing as well (referential containment). If you want them unfiled from `/labs/User-02` to complete the move, just say the word.

New task

- Perfect, it’s all done!
- OK, let’s review the properties of one document from the list:

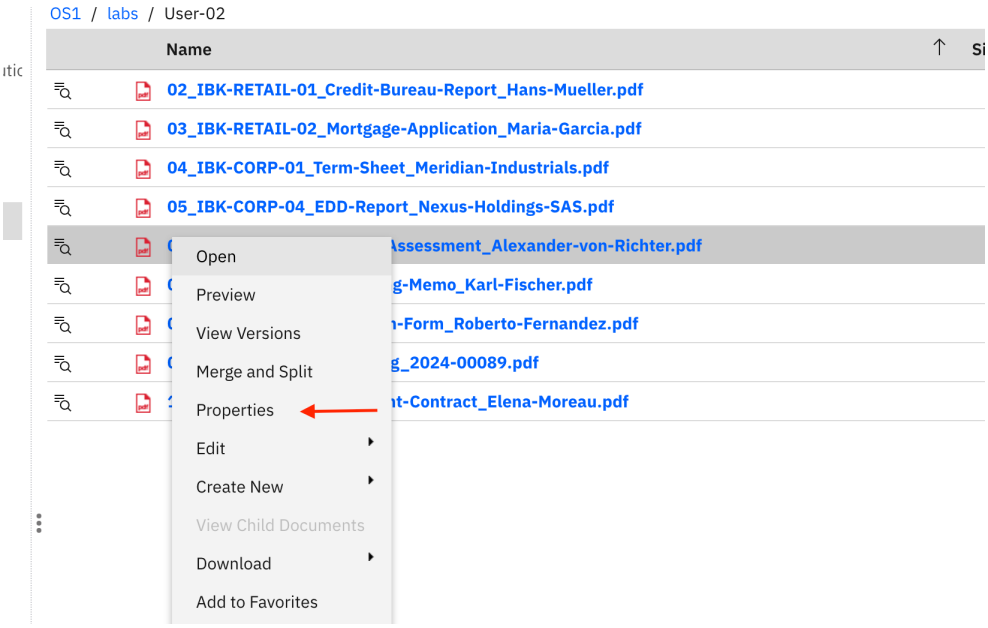

- Bob has updated the properties of the document:

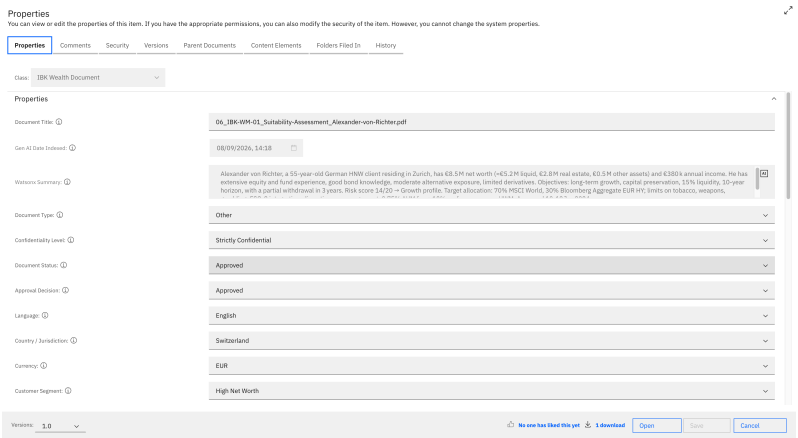

- Bob has added the document to the right folder. You can confirm it by verifying the “Folder Field In” in the properties page of the document :

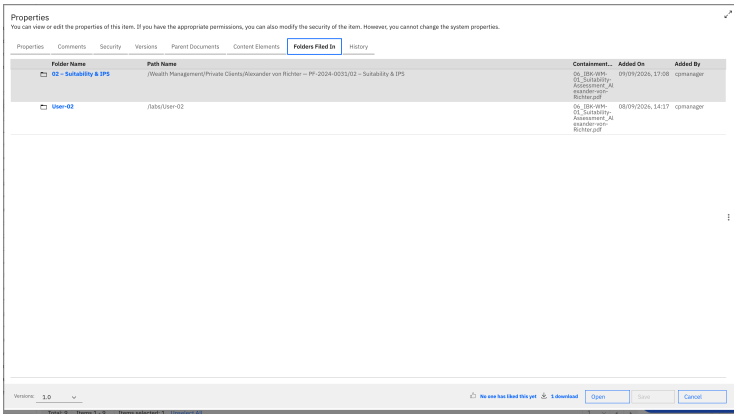

5.2 Content Cortex Assistant Integration (Optional) :

https://www.ibm.com/docs/en/content-assistant?topic=developing-generative-ai-applications

Objective: Explore how developers can embed governed content intelligence in custom applications. You will use GraphQL for repository and generative-AI operations, then relate the exercise to the broader Content Cortex integration model, where MCP-based services can expose authorised content tools to AI agents.

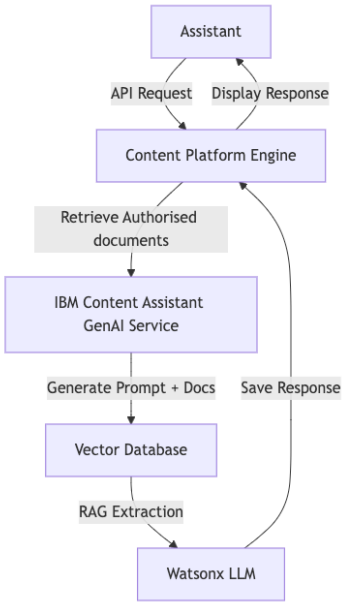

5.2.1 Graphql

GraphQL is a query-language specification; you can read more about it at https://graphql.org/. It was designed to address limitations in RESTful interfaces and, like those interfaces, uses HTTP. A key advantage of GraphQL is its flexibility in defining the information returned in a response. As a result, several requests that would be separate in traditional FileNet Content Platform Engine interfaces can be combined into one request for greater efficiency.

Another advantage is that applications interfacing with GraphQL only require sending and receiving HTTP messages and can be written without binding with any specific FileNet Content Platform Engine Client libraries. This results in fewer version dependencies and easier environment upgrades.

For further reading, the FileNet P8 Platform Documentation contains some sections about developing with the GraphQL API for Content Platform Engine:

https://www.ibm.com/docs/en/filenet-p8-platform/5.6.0?topic=development-overview

Furthermore, there is also a white paper on the GraphQL interface, with lots of examples: https://www.ibm.com/support/pages/sites/default/files/inline-files/$FILE/IBM%20Content%20Services%20GraphQL%20API%20Developer%20Guide_1.pdf.

For further examples:

https://www.ibm.com/docs/en/content-assistant?topic=applications-content-engine-java-api-samples

https://www.ibm.com/docs/en/content-assistant?topic=applications-content-engine-graphql-api-sample-mutations-queries

5.2.2 Lab Setup Instructions

If you are performing this lab as part of an IBM event, access the document that lists the available systems, URLs, and login instructions. For this lab, you will need to access **Graph*i*QL**.

- Paste the Content Services GraphiQL URL into your web browser.

  https://cpd-cp4ba.apps.xxx.eu1.techzone.ibm.com/content-services-graphql

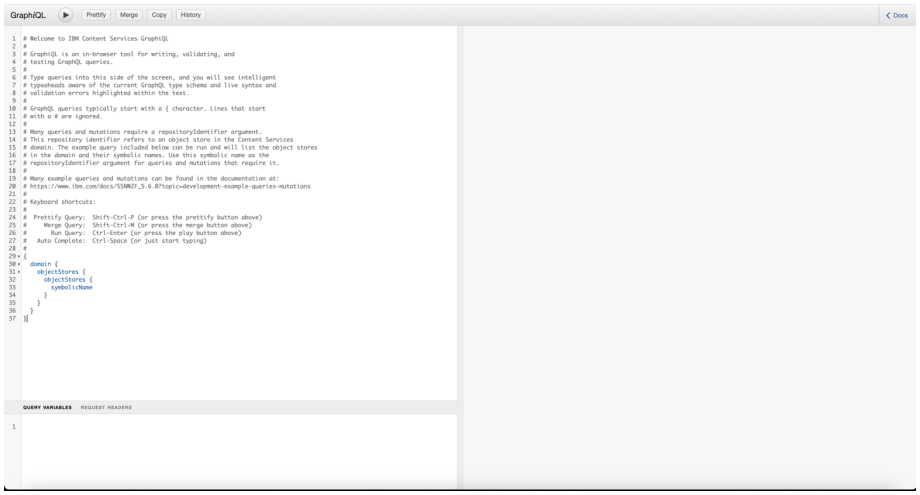

- You hit the play button to test the Graph*i*QL

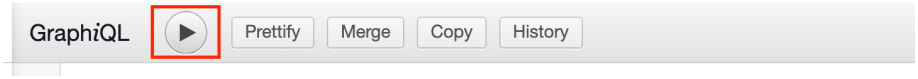

The result should look like:

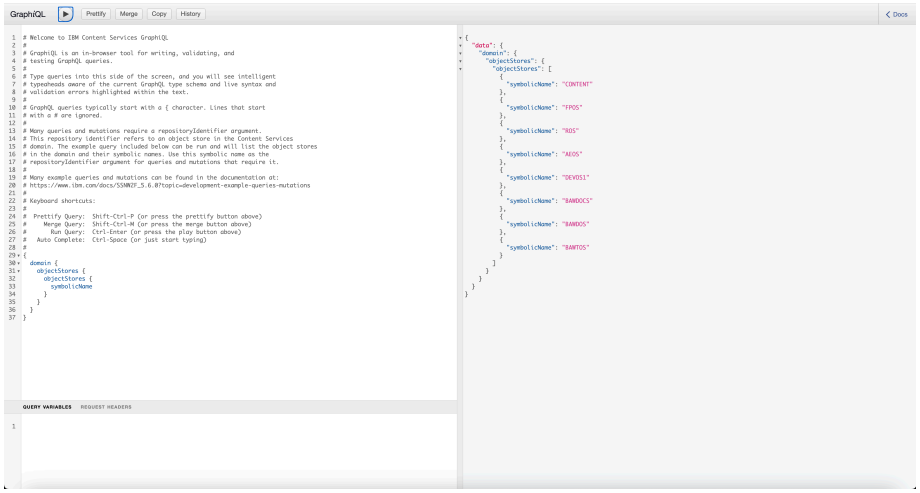

Perfect, let's do some Gen AI now.

Let's open the following page: https://github.com/MalekJabri/ContentAssistant/blob/main/lab-01/readme.md

In this lab, you will copy code snippets from the GitHub page into the corresponding sections of the GraphiQL interface. Be sure to paste each snippet into the correct section.

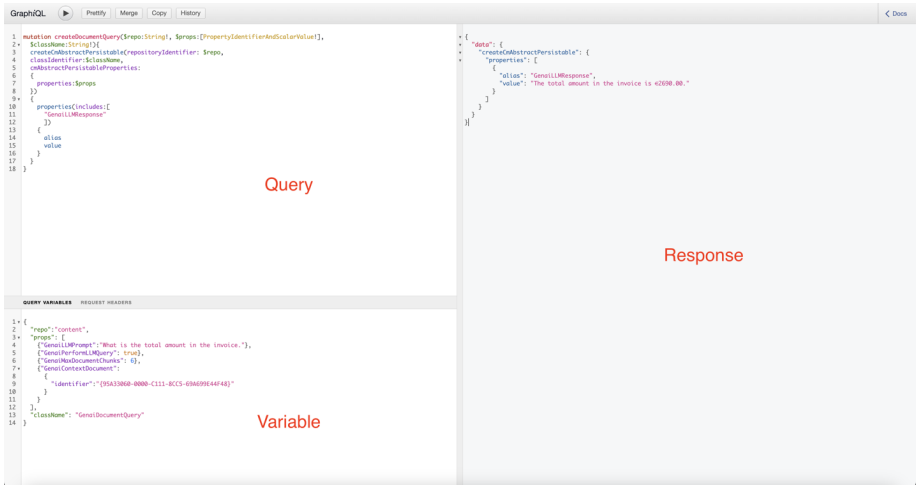

5.2.3 Ask a generic question.

Our initial test will involve asking a generic question utilising the embedded language model (LLM), which is unrelated to any specific document in the repository.

Copy of the query in the first section

Step 1: Ask a generic question

Copy and paste on to each section :

Query section

```
mutation createBaseQuery($repo:String!, $props:[PropertyIdentifierAndScalarValue!],
  $className:String!){
  createCmAbstractPersistable(repositoryIdentifier: $repo,
  classIdentifier:$className,
  cmAbstractPersistableProperties:
  {
    properties:$props
  })
  {
    id
    name
    creator
    properties(includes:["GenaiLLMPrompt",
                         "GenaiLLMModelName",
                            "GenaiLLMResponse",
                         ])
    {
      alias
      value
    }
  }
}
```

Copy the first example from the variable section.

"GenaiLLMPrompt": The prompt is the text you provide to guide the system or model's response. It sets the context or topic for the task you're asking it to perform.

Commented [GV1]: This is unclear. Direction should specifically state they should copy the first example for the Variable section in the Github page and paste into the Query Variable field.

```
{
  "repo":"content",
  "props": [
    {"GenaiLLMPrompt":"Describe ECM"}
  ],
  "className": "GenAIBaseQuery"
}
```

You can click on the play button

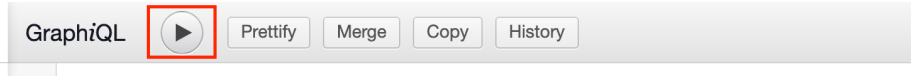

You can review the response. The answer is in the properties section for the value of the field "GenAILLMResponse".

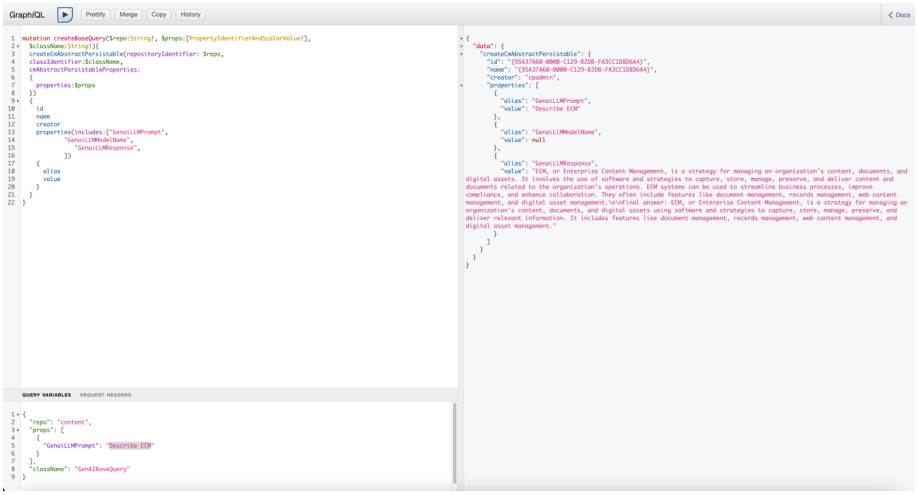
GraphiQL
Prettify
Merge
Copy
History
Docs
mutation createBaseQuery($repo:String!, $props:[PropertyIdentifierAndScalarValue!],
$classNam:String){
createCmAbstractPersistable(repositoryIdentifier: $repo,
classIdentifier:$classNam,
cmAbstractPersistableProperties:
{
properties:$props
})
{
id
name
creator
properties(includes:["GenaiLLMPrompt",
"GenaiLLMModelName",
"GenaiLLMResponse",
])
{
alias
value
}
}
}
{
"data": {
"createCmAbstractPersistable": {
"id": "{95A37A60-0000-C129-B2D8-FA3CC1D8D6A4}",
"name": "{95A37A60-0000-C129-B2D8-FA3CC1D8D6A4}",
"creator": "cpadmin",
"properties": [
{
"alias": "GenaiLLMPrompt",
"value": "Describe ECM"
},
{
"alias": "GenaiLLMModelName",
"value": null
},
{
"alias": "GenaiLLMResponse",
"value": "ECM, or Enterprise Content Management, is a strategy for managing an organization's content, documents, and digital assets. It involves the use of software and strategies to capture, store, manage, preserve, and deliver content and documents related to the organization's operations. ECM systems can be used to streamline business processes, improve compliance, and enhance collaboration. They often include features like document management, records management, web content management, and digital asset management.\n\nFinal answer: ECM, or Enterprise Content Management, is a strategy for managing an organization's content, documents, and digital assets using software and strategies to capture, store, manage, preserve, and deliver relevant information. It includes features like document management, records management, web content management, and digital asset management."
}
]
}
}
}
QUERY VARIABLES
REQUEST HEADERS
{
"repo": "content",
"props": [
{
"GenaiLLMPrompt": "Describe ECM"
}
],
"className": "GenAIBaseQuery"
}

5.2.4 Ask your repository

In this exercise, you will leverage the "Ask Your Repository" feature to efficiently submit questions and access comprehensive information from the entire documentation repository using the GraphQL API

Copy of the query in the first section

Step 2: Ask your repository

Query Section

```
mutation createVectorQuery($repo:String!, $props:[PropertyIdentifierAndScalarValue!],
  $className:String!){
  createCmAbstractPersistable(repositoryIdentifier: $repo,
  classIdentifier:$className,
  cmAbstractPersistableProperties:
  {
    properties:$props
  })
  {
    id
    name
    creator
    properties(includes:["GenaiLLMPrompt",
      "GenaiLLMModelName",
      "GenaiLLMResponse",
      "GenaiVectorChunks"
      "GenaiPerformLLMQuery"
      "GenaiMaxDocumentChunks"
      ])
    {
      alias
      value
    }
  }
}
```

Copy the first example from the variable section.

“GenAIMaxDocumentChunks”: The number of chunks used to generate the answer. You can change the number of chunks to provide more context.

Variable Section

```
{
  "repo": "content",
  "props": [
    {
      "GenaiLLMPrompt": "Does AXA home policy cover Vehicles or craft?"
    },
    {
      "GenaiPerformLLMQuery": true
    },
    {
      "GenaiMaxDocumentChunks": 6
    }
  ],
  "className": "GenaiVectorQuery"
}
```

You can click on the play button.

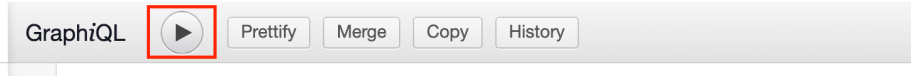

You can review the response. The answer is in the properties section for the field value “GenAILLMResponse”.

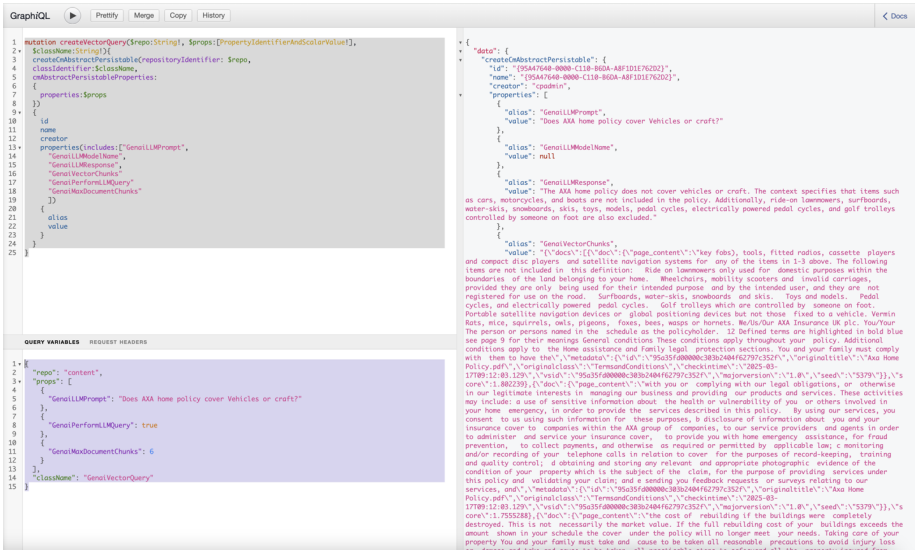

You can also review the various segments used to generate the answer in the properties “GenAIVectorChunks”.

5.2.5 Ask a question on a document.

This exercise will use Graphql to ask questions about a specific document in our repository. By crafting targeted queries, we can access detailed insights directly related to the chosen document.

We will first list a set of documents in a specific folder to retrieve their IDS.

Copy of the query in the section

Step 3: list documents inside a folder.

Query Section

```
{
  folder(repositoryIdentifier: "content", identifier: "/Focus Corp/Invoices") {
    containedDocuments {
      documents {
        id
        name
      }
    }
  }
}
```

Retrieve the document list and copy the results into a text editor. Use one of the IDs in the following exercise.

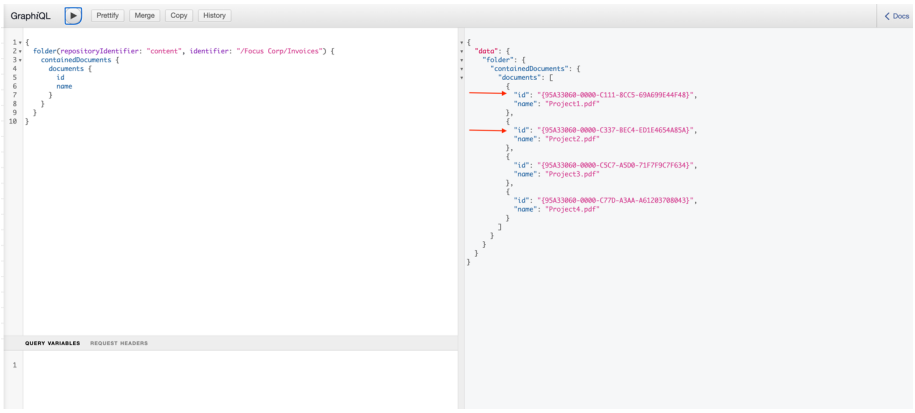

Copy the text in the query section

Commented [GV2]: This should after "(Step 4) Gen AI on specific document.

(Step 4) Gen AI on specifc document

```
mutation createDocumentQuery($repo:String!, $props:[PropertyIdentifierAndScalarValue!],
  $className:String!){
  createCmAbstractPersistable(repositoryIdentifier: $repo,
  classIdentifier:$className,
  cmAbstractPersistableProperties:
  {
    properties:$props
  })
  {
    id
    name
    creator
    properties(includes:["GenaiLLMPrompt",
      "GenaiLLMModelName",
      "GenaiLLMResponse",
      "GenaiVectorChunks"
      "GenaiPerformLLMQuery"
      "GenaiMaxDocumentChunks"
      "GenaiContextDocument"
      ])
    {
      alias
      value
    }
  }
}
```

Copy the next part to the variable section.

variable section

```
{
  "repo":"content",
  "props": [
    {"GenaiLLMPrompt":"What is the total amount in the invoice."},
    {"GenaiPerformLLMQuery": true},
    {"GenaiMaxDocumentChunks": 6},
    {"GenaiContextDocument":
        {
      → "identifier":"{95A33060-0000-C111-8CC5-69A699E44F48}"
      }
    }
  ],
  "className": "GenaiDocumentQuery"
}
```

Update the identifier using one of the IDs retrieved in the previous step.

You can click on the play button.

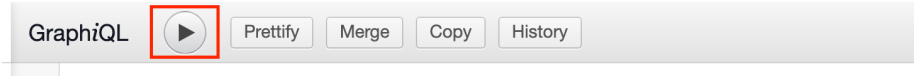

You can review the response. The answer is in the properties section for the field value “GenAILLMResponse”.

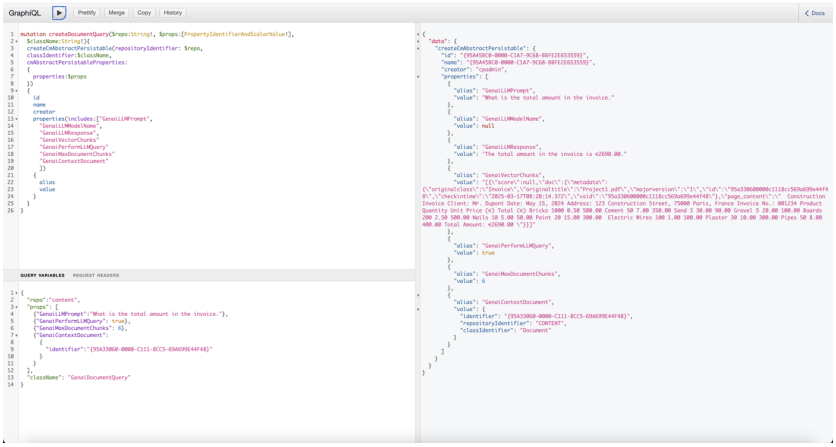

5.2.6 More exercises:

In the GitHub page, we are also exploring additional examples like:

Gen AI is on a single document with a simplified version.

Gen AI on multiple documents

Search for historical requests.

6 Lab Exercise 5: Configure and govern AI-enabled content (Optional)

Objective: Review the controls that shape AI behaviour and content eligibility. Content Assistant configuration is one part of the wider Content Cortex governance model, which also includes repository security, metadata, lifecycle controls, auditability, and controlled exposure of content operations to agents.

In this section, you'll explore key configuration options, including:

- **Setting custom prompts** – Define how AI summarises documents and responds to queries.
- **Selecting the AI model** – Choose the best model based on performance and use case.
- **Tracking past queries** – Review and analyse previously asked questions to refine usage.
- **Activating document classes** – Control which types of documents are enabled for AI-powered assistance.

6.1 Introduction

Behind the scenes, IBM Content Assistant relies on key **backend and frontend components** to ensure smooth operation. This section overviews the essential administrative elements supporting AI-driven document assistance.

Here's what you'll learn:

1. **Plugins & Extensions** – These components are required for backend and frontend functionality.
2. **Object Classes & Properties** – How different document structures and metadata are imported by the system.
3. **Indexing & Automation** – Understanding key actions, including **sweep jobs**, to process and index existing content.

By the end of this section, you'll have a solid understanding of **how IBM Content Assistant is structured** and what's needed to manage it effectively.

6.2 Enable your content assistant and samples (Front)

IBM Content Navigator acts as the **front-end framework** for engaging with enterprise content. It enables administrators to **install plugins** that introduce new features and enhance user interface capabilities. A key aspect of customisation is creating desktops that operate like tailored **applications**.

Desktops allow organisations to specify **who has access to certain features and content**, customising the user experience according to roles and needs.

1. Connect to the Admin Section of Content Navigator and click on the “hamburger menu” :

https://cpd-cp4ba.apps.xxxx.eu1.techzone.ibm.com/icn/navigator/?desktop=CONTENT

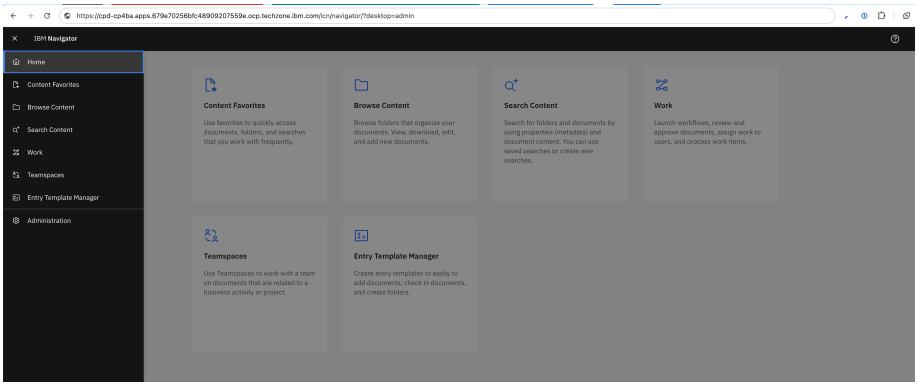

2. Click on the Administration option at the bottom of the left panel.

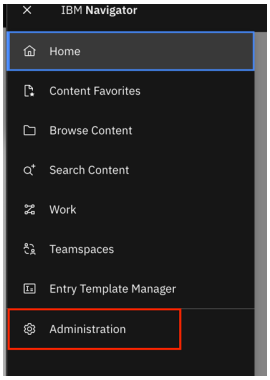

3. Create a new desktop and select the “Platform and Content” option.

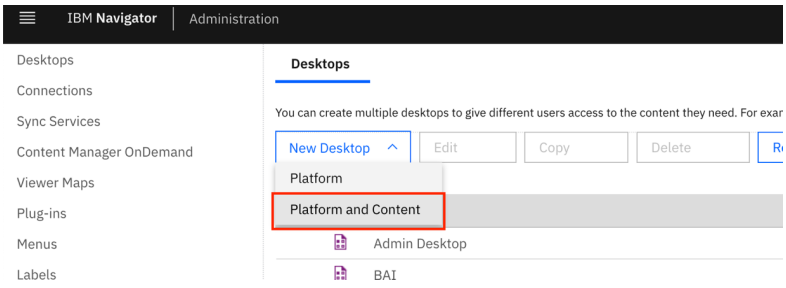

4. Set the name as "USXX Legal Assistant", where XX is the unique number given to you by the instructor.
5. Select Content as Repository

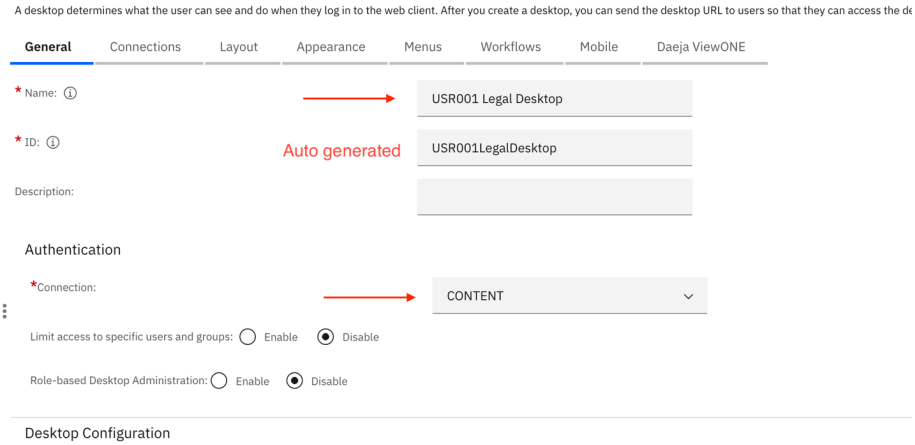

6. Collapse the Desktop Configuration section and validate that "enable all plug-ins..." are selected in the plugin section.

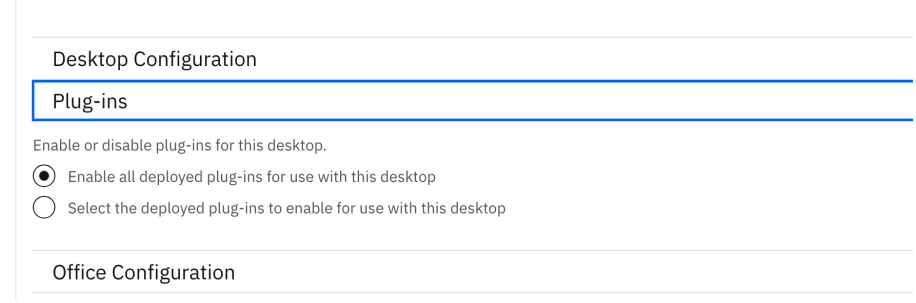

7. Click on the tab layout and select “Ask Your Repository” on the Displayed features list:

Commented [RE3]: It is called Ask Your Repository on the system I have access to.

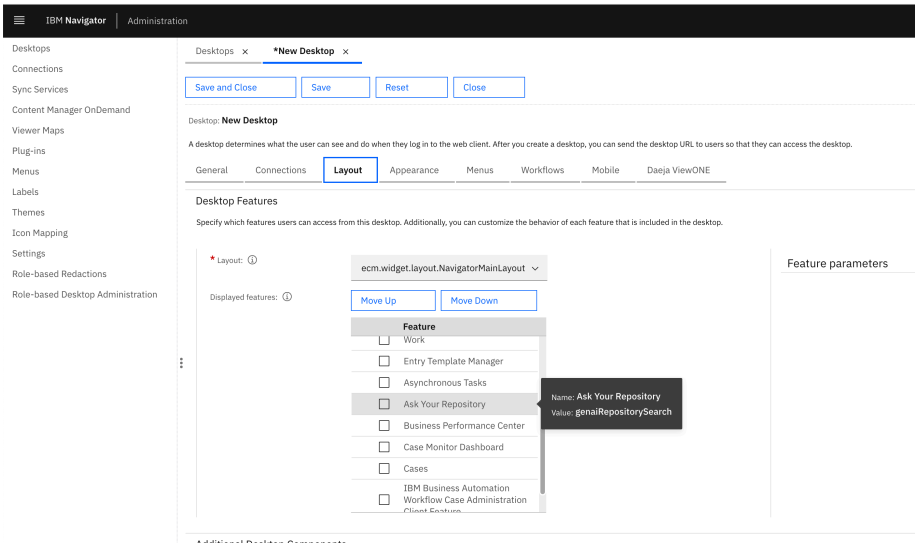

8. Click on “Save and Close”

Save and Close | Save | Reset | Close

Desktop: **New Desktop**

9. Select the plugin option.

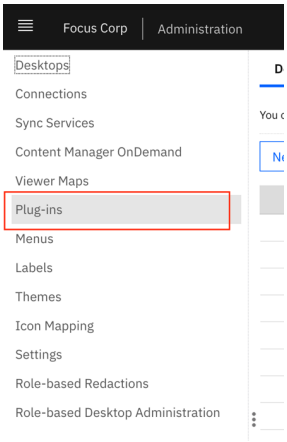

10. Select the Content Assistant Plugins and click on the Edit button

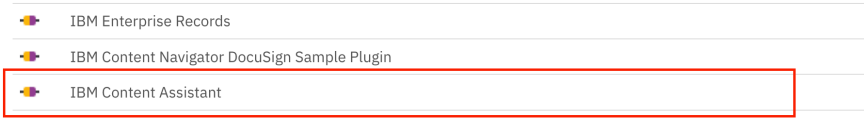

11. You can customise your predefined questions for our new desktop.

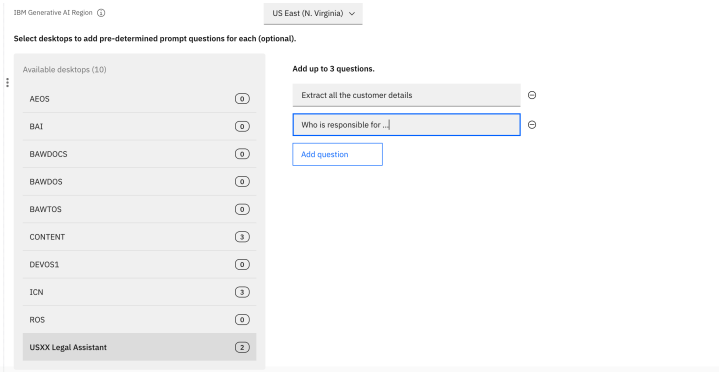

12. Navigate to the new desktop by replacing admin with USXXLegalAssistant.

m/navigator/?desktop=USXXLegalAssistant

6.3 Administration and operational assurance

6.3.1 Review Content Cortex backbone configuration

1. Navigate to the Administrative Console for Content Platform Engine. The link ends with "/CPE/acce/."

   https://xxxxxxx.eu1.techzone.ibm.com/cpe/acce/

2. Select the Object Stores called "OS1." An object store is a library where content is classified, stored, and secured.

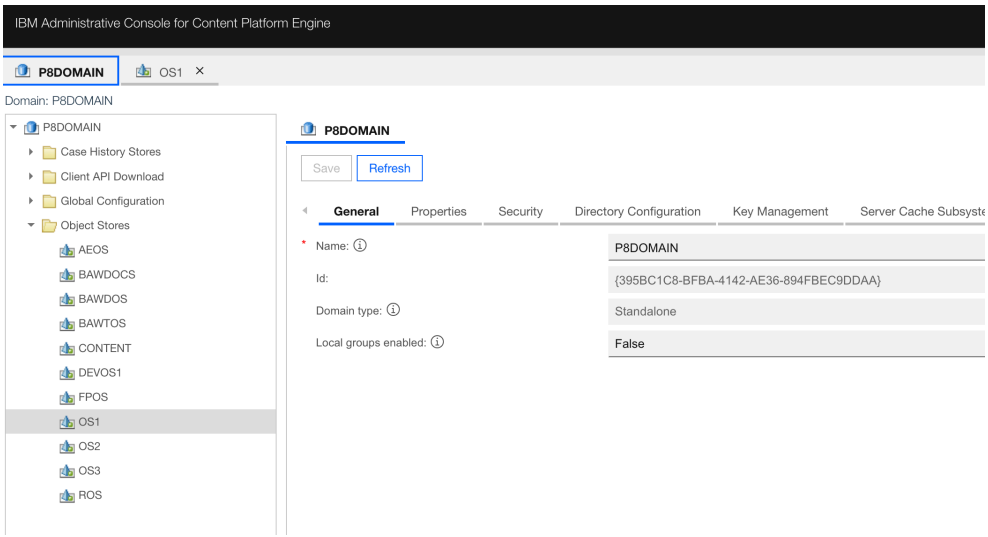

3. A new tab, “OS1”, will be open. That will provide you access to all the configurations possible for that Object Store. Click on the Search option on the left panel.

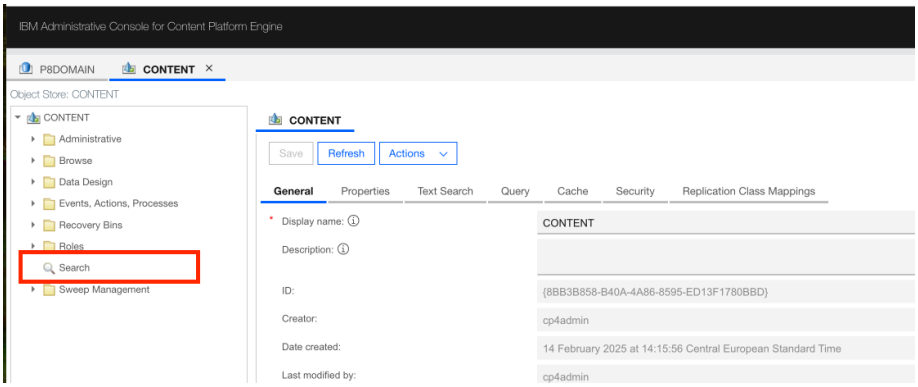

- Since version 26, we have introduced a Gen AI configuration section in ACCE

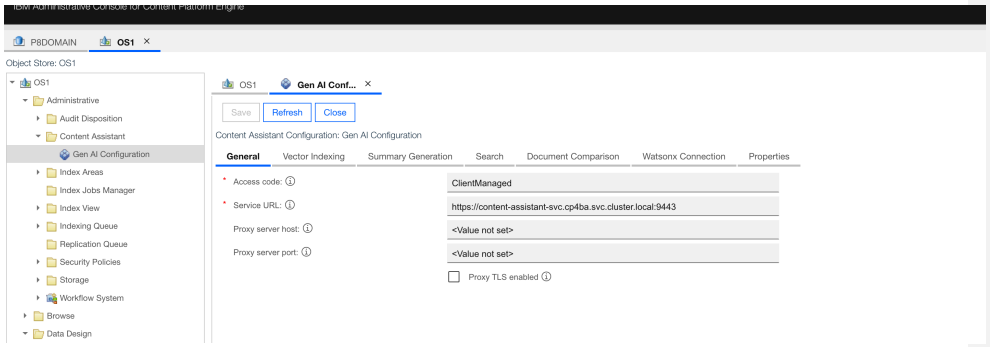

4. Please review the Gen AI Configuration

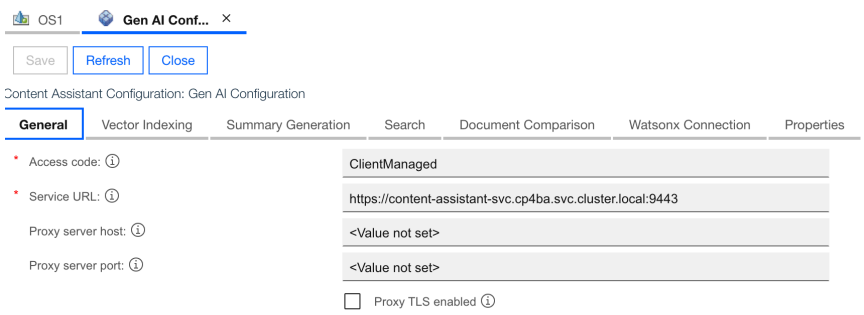

5. This object contains the major configurations for the Content Assistant. DO NOT CHANGE ANY OF THESE SETTINGS!

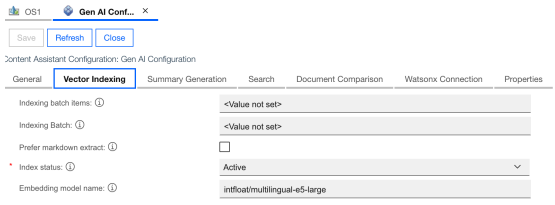

6. Click on the button next to the Summary Generation tab to review the content.

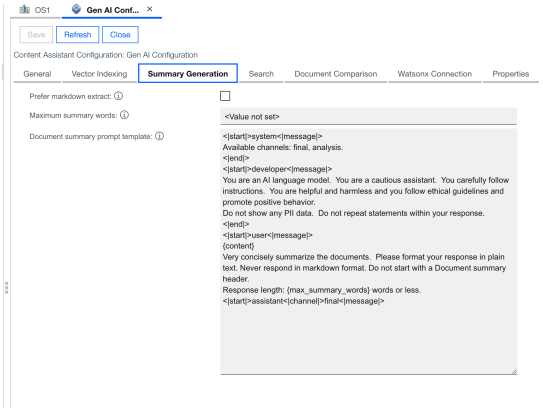

7. You can review the default prompt on search tab. You can also customise the prompt to fit your business requirements.

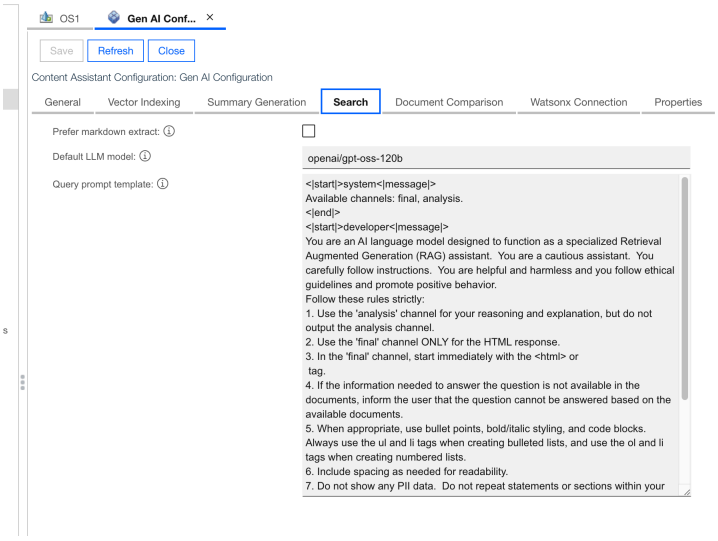

6.3.2 Enhance Extraction (New feature)

The Content Engine includes text extraction capabilities for document indexing and content-based retrieval. By default, it extracts and indexes text but doesn't save it. With persistent text extraction, the extracted text is stored as an annotation on each document through the Text Extraction Annotation class. This stored text supports multiple purposes: content-based retrieval (CBR) indexing, vector indexing, document summarization, and generative AI inferences. For more information about the persistent text extraction classes and properties, see Persistent Text Extract Extensions.

Enhanced extraction builds on persistent text extraction by introducing an extensible handler framework. This framework allows you to integrate specialised text extraction engines alongside the default Oracle Outside In Technology (OIT) engine. Each specialised handler can extract additional metadata and process complex document types that OIT cannot handle effectively.

We have introduced a set of new enhanced extraction mechanisms:

- **Setting up PDF Forms enhanced extraction handler**
  Set up the PDF Forms enhanced extraction handler to extract form fields from PDF documents with proper field-to-value associations.
- **Setting up watsonx.ai enhanced extraction handler**
  Extract text from images and structured data by configuring the watsonx.ai enhanced extraction handler.

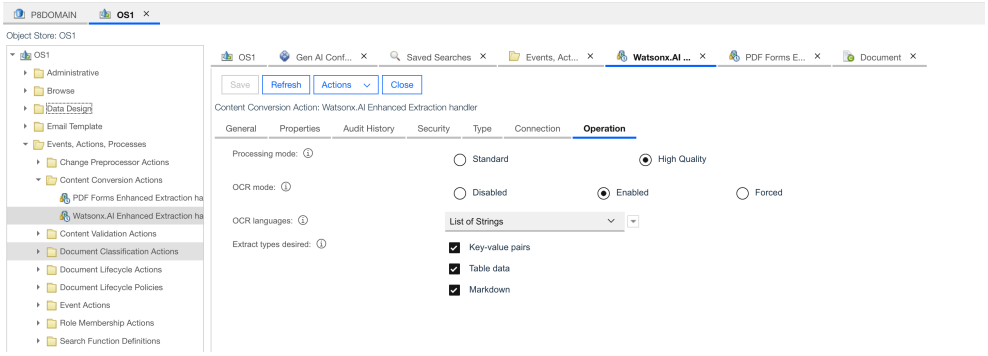

6.3.3 Monitor AI requests and inspect traceability

In this section, we can monitor all the Gen AI requests made in the platform. This includes the following elements.

All documents that will be indexed in the vector database

All document summary requests

All questions related to a single document

All questions related to multiple documents.

1. In the Administrative Console for Content Platform Engine, select the search Option on the left panel.

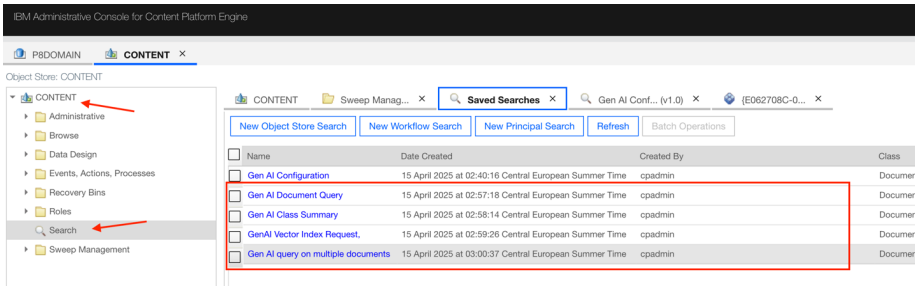

2. There are four predefined searches.
3. In the list, select the “Gen AI Query On multiple documents” search. Then click Run and OK in the confirmation box. This search

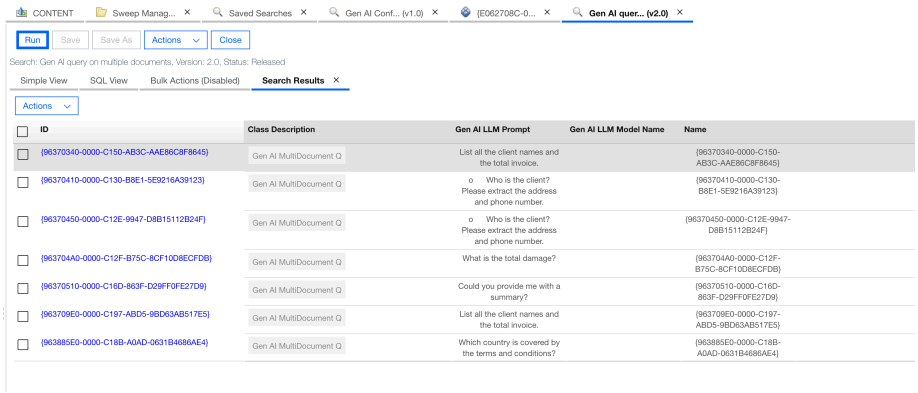

4. You can review all the questions different users ask, including the responses and the chunks used to provide the answers.
5. Select any entry in the list by clicking on the ID.

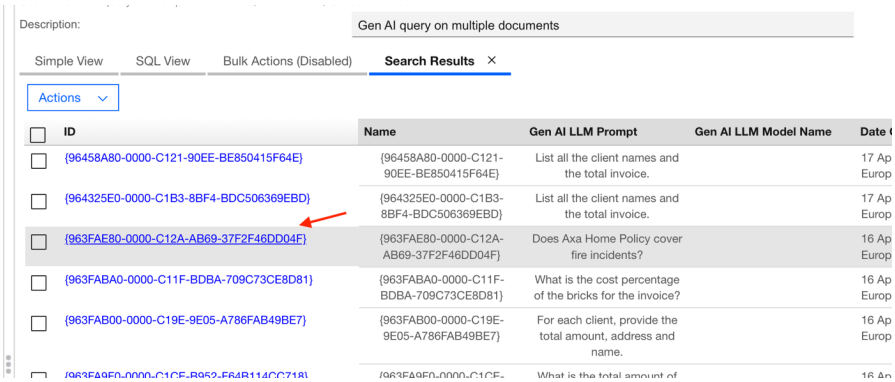

6. We will be able to review the context of the question, including

    1. Creator: The user who asks the question
    2. Gen AI Context Documents: The list of IDS of documents used to answer the question
    3. Gen AI LLM Prompt: The question that has been asked.
    4. Gen AI LLM Response: The response provided
    5. Gen AI Max Document Chunks: The maximum number of extractions.
    6. Gen AI Vector Chunks: The extracted text generated the answer, including the score, ID document, title, ….

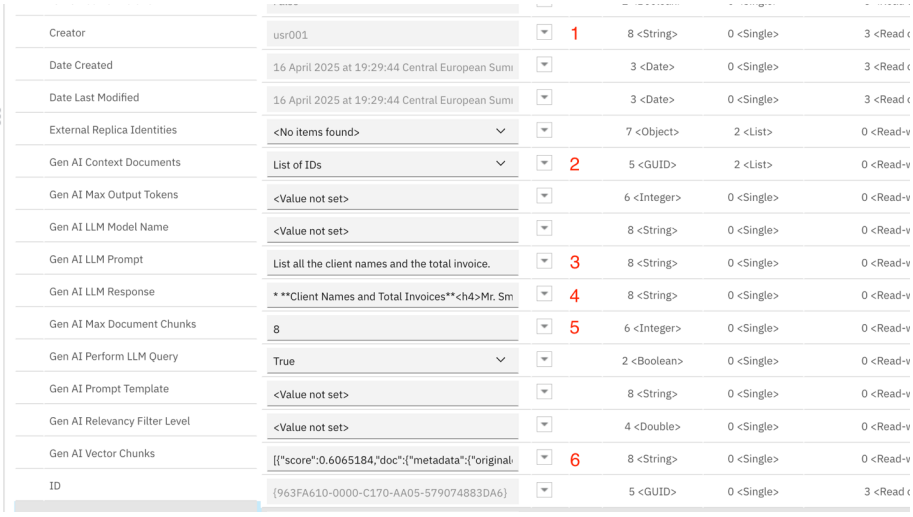

7. For example, to review the details, click the arrow next to the value to review the answer.

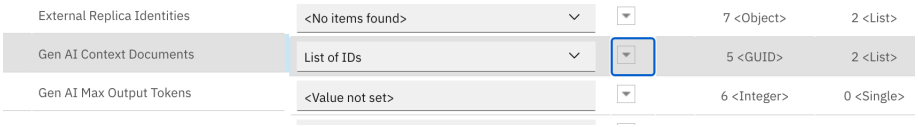

8. Click on display and edit value:

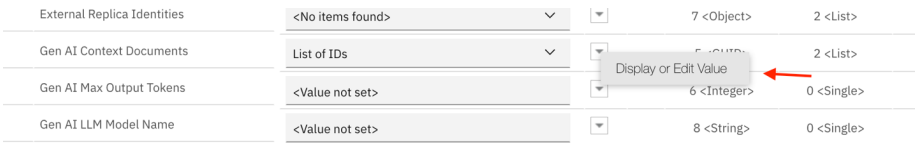

9. This represents the list of documents ID: